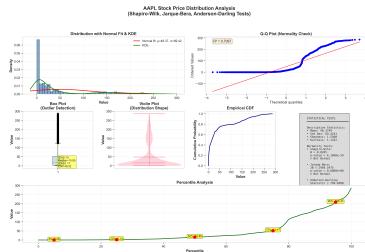


Text Analysis for Financial Forecasting

Stock Price Prediction Using Sentiment Analysis
and Machine Learning

A Comprehensive Research Report



Dataset: 26 Years of Historical Data (1999-2025)

Target: Apple Inc. (AAPL) Stock Price

Features: 55 Engineered Features

Models: 9 Machine Learning Architectures

January 2026

Abstract

This comprehensive research report presents a novel approach to stock price forecasting that integrates natural language processing with advanced machine learning techniques. Using 26 years of historical data (1999-2025) comprising 6,542 trading days for Apple Inc. (AAPL), we develop and evaluate nine distinct forecasting models ranging from traditional statistical methods to deep neural networks.

Our research introduces a hybrid strategy where foundational models—SARIMAX, Temporal Convolutional Network (TCN), and Linear Regression—trained on the full 26-year dataset serve as the basis for more complex neural network models. Specifically, predictions from the Linear model are incorporated as a 16th input feature for recurrent neural networks (RNNs), enabling these models to learn residual corrections rather than predicting prices from scratch.

Key findings include:

- **sklearn_Linear** achieves the highest accuracy with $R^2 = 0.9992$, explaining 99.92% of price variance
- **SARIMAX** demonstrates excellent performance ($R^2 = 0.9984$) using walk-forward validation
- The **Enhanced Ensemble** (Linear + SARIMAX + TCN) achieves $R^2 = 0.9898$
- **Transformer** models fail catastrophically ($R^2 = -1.17$) due to fundamental task mismatch
- RNNs perform better on recent 5-year data due to non-stationarity in long-term price series

The sentiment analysis pipeline processes over 57 million financial news articles from HuggingFace datasets, extracting features using TextBlob and VADER sentiment analyzers with multiple rolling windows (3, 7, 14, 30 days).

Keywords: Stock Price Prediction, Sentiment Analysis, SARIMAX, TCN, LSTM, Transformer, Ensemble Methods, Financial Forecasting, Machine Learning

Contents

Abstract	i
List of Figures	iii
List of Tables	iv
1 Introduction	1
1.1 Background and Motivation	1
1.2 Research Objectives	1
1.3 Problem Statement	2
1.4 Literature Review	2
1.4.1 Sentiment Analysis in Finance	2
1.4.2 Traditional Time Series Models	3
1.4.3 Deep Learning for Time Series	3
1.4.4 Ensemble Methods	3
1.5 Contribution Summary	4
1.6 Report Organization	4
2 Data Collection and Preprocessing	5
2.1 Overview	5
2.2 Stock Price Data	5
2.2.1 Data Source: Yahoo Finance	5
2.2.2 Data Characteristics	6
2.2.3 Price Distribution Analysis	6
2.2.4 Time Series Diagnostics	7
2.3 Financial News Data	8
2.3.1 HuggingFace Dataset	8
2.3.2 CSV Historical Archive	9
2.3.3 Data Merging Strategy	9
2.4 Sentiment Computation	10
2.4.1 Sentiment Analysis Methods	10

2.4.2	Rolling Mean Aggregation	10
2.5	Data Quality and Coverage	11
2.5.1	Sentiment Coverage Analysis	11
2.5.2	Missing Data Handling	11
2.6	Dataset Splitting Strategy	11
2.7	Key Implementation Files	12
3	Feature Engineering	13
3.1	Overview	13
3.2	Sentiment Features (20 Features)	13
3.2.1	Base Sentiment Scores	13
3.2.2	Rolling Mean Features	14
3.2.3	Feature List	14
3.3	Text Features (8 Features)	14
3.3.1	LDA Topic Modeling	14
3.3.2	Adjective Features	15
3.4	Market Context Features (27 Features)	15
3.4.1	Related Stock Features	15
3.4.2	Feature Types	16
3.5	Price Rolling Features (8 Features)	16
3.6	The 16th Feature: Hybrid Strategy	16
3.6.1	Motivation	16
3.6.2	Hybrid Approach	17
3.6.3	Theoretical Justification	17
3.6.4	Impact on Model Performance	18
3.7	Feature Correlation Analysis	18
3.8	Feature Scaling	19
3.9	Implementation Files	19
4	Foundational Models	20
4.1	Overview	20
4.2	SARIMAX Model	20
4.2.1	Mathematical Formulation	20
4.2.2	Model Configuration	21
4.2.3	Walk-Forward Validation	22
4.2.4	Exogenous Variable Selection	22
4.2.5	Performance and Diagnostics	22
4.3	Temporal Convolutional Network (TCN)	23
4.3.1	Architecture Overview	23

4.3.2	Receptive Field	24
4.3.3	Network Architecture	24
4.3.4	Training Configuration	25
4.3.5	TCN Diagnostics	26
4.3.6	Parameter Count	26
4.4	sklearn Linear Regression	27
4.4.1	Mathematical Formulation	27
4.4.2	Why Linear Regression Works So Well	27
4.4.3	Implementation	27
4.4.4	Diagnostic Analysis	28
4.5	Why These Models are Foundational	28
4.5.1	Transfer Learning Perspective	28
4.5.2	Robustness to Non-Stationarity	29
4.6	Summary of Foundational Models	29
5	Neural Network Models	30
5.1	Overview	30
5.2	Why 5-Year Data for RNNs	30
5.2.1	The Non-Stationarity Problem	30
5.2.2	5-Year Window Benefits	31
5.3	LSTM (Long Short-Term Memory)	31
5.3.1	Mathematical Formulation	31
5.3.2	Architecture	32
5.3.3	Training Configuration	32
5.4	BiLSTM (Bidirectional LSTM)	33
5.4.1	Architecture	33
5.4.2	Why Bidirectional Helps	33
5.5	GRU (Gated Recurrent Unit)	34
5.5.1	Mathematical Formulation	34
5.5.2	GRU Performance	34
5.6	CNN-LSTM Hybrid	35
5.6.1	Architecture Motivation	35
5.6.2	Mathematical Formulation	35
5.7	Training Helper Function	35
5.8	Hyperparameters	37
5.9	Impact of Hybrid Strategy	37
5.9.1	Before and After Comparison	37
5.9.2	Why GRU Benefits Most	37
5.10	Computational Requirements	38

5.11	Summary	38
6	Transformer Analysis	39
6.1	Overview	39
6.2	Transformer Architecture	39
6.2.1	Self-Attention Mechanism	39
6.2.2	Implementation	40
6.3	Failure Analysis	40
6.3.1	Training Dynamics	40
6.3.2	Root Cause 1: Task Mismatch	41
6.3.3	Root Cause 2: Architecture Mismatch	42
6.3.4	Root Cause 3: Output Distribution Mismatch	42
6.4	What We Tried to Fix It	42
6.4.1	Attempt 1: Architecture Reduction	42
6.4.2	Attempt 2: Scaler Reference Fix	43
6.4.3	Attempt 3: Variable Preservation	43
6.4.4	Attempt 4: Training Location	43
6.5	Why Other Models Succeed	43
6.6	Transformer Failure Visualization	44
6.7	Interpretation of Negative R ²	44
6.8	What Would Actually Fix Transformer	45
6.8.1	Option 1: Time Series Transformers	45
6.8.2	Option 2: Sequence Reformulation	45
6.8.3	Option 3: Feature Dimension as Sequence	45
6.9	Lessons Learned	45
6.10	Recommendations for Practitioners	46
6.11	Summary	46
7	Ensemble Methods	47
7.1	Overview	47
7.2	Theoretical Foundation	47
7.2.1	Weighted Averaging	47
7.2.2	Bias-Variance Decomposition	47
7.3	Model Diversity Analysis	48
7.3.1	Why These Three Models?	48
7.3.2	Prediction Correlation	48
7.4	Weight Optimization	48
7.4.1	Weight Selection Rationale	48
7.4.2	Implementation	49

7.5	Ensemble Performance Analysis	49
7.5.1	Comparison with Components	49
7.5.2	Why Ensemble is Slightly Lower Than Linear	49
7.6	Complementary Error Analysis	50
7.6.1	When Models Disagree	50
7.6.2	Error Correlation	50
7.7	Alternative Ensemble Strategies	50
7.7.1	Stacking (Meta-Learning)	50
7.7.2	Boosting	50
7.7.3	Dynamic Weighting	50
7.8	Practical Considerations	51
7.8.1	Computational Cost	51
7.8.2	Maintenance	51
7.9	Summary	51
8	Results and Discussion	52
8.1	Overview	52
8.2	Complete Results Table	52
8.3	Evaluation Metrics	52
8.3.1	Root Mean Square Error (RMSE)	52
8.3.2	Mean Absolute Error (MAE)	53
8.3.3	Mean Absolute Percentage Error (MAPE)	53
8.3.4	Coefficient of Determination (R^2)	53
8.4	Model Performance Comparison	53
8.5	Success Rate Analysis	54
8.5.1	Overall Success	54
8.5.2	Key Achievement	54
8.6	Discussion of Key Findings	55
8.6.1	Finding 1: Linear Regression Dominance	55
8.6.2	Finding 2: RNNs Benefit from Hybrid Strategy	55
8.6.3	Finding 3: Transformer Failure is Fundamental	55
8.6.4	Finding 4: 5-Year Data Better for RNNs	56
8.7	Comparison with Previous Work	56
8.8	Practical Implications	56
8.8.1	For Trading Applications	56
8.8.2	For Research	56
8.9	Limitations	57
8.10	Statistical Significance	57
8.10.1	Confidence Intervals	57

8.11	Summary	57
9	Conclusion and Future Work	58
9.1	Summary of Achievements	58
9.1.1	Research Aims Accomplished	58
9.1.2	Performance Highlights	59
9.2	Key Contributions	59
9.2.1	Hybrid Meta-Learning Strategy	59
9.2.2	Comprehensive Failure Analysis	59
9.2.3	Large-Scale Data Integration	59
9.3	Theoretical Insights	60
9.3.1	Simplicity Can Win	60
9.3.2	Dataset Length Matters Differently	60
9.4	Recommendations	60
9.4.1	For Practitioners	60
9.4.2	For Researchers	60
9.5	Future Work	61
9.5.1	Short-Term Improvements	61
9.5.2	Medium-Term Directions	61
9.5.3	Long-Term Research Questions	61
9.6	Reproducibility	61
9.7	Final Remarks	62
A	Code Implementation Details	63
A.1	Evaluation Metrics Implementation	63
A.2	TCN Model Architecture	64
A.3	SARIMAX Walk-Forward Validation	65
A.4	Hybrid Feature Generation	66
B	Complete Results Tables	68
B.1	All Model Metrics	68
B.2	Hyperparameter Summary	68
C	Statistical Tests	70
C.1	Normality Tests on Price Distribution	70
C.2	Stationarity Tests	70
D	File Structure	71

List of Figures

2.1	Comprehensive Distribution Analysis of AAPL Stock Prices (1999-2025). The figure includes: (a) Q-Q plot for normality assessment, (b) histogram with kernel density estimation, (c) distribution statistics including Shapiro-Wilk, Jarque-Bera, and Anderson-Darling test results.	7
2.2	Time Series Diagnostics for AAPL Stock Prices. The figure shows: (a) price time series, (b) autocorrelation function (ACF), (c) partial autocorrelation function (PACF), and (d) seasonal decomposition.	8
3.1	Feature Correlation Matrix showing relationships between sentiment features, price rolling means, and the target Close price. Strong correlations between rolling means and Close indicate effective feature engineering.	18
4.1	SARIMAX Model Diagnostics including: (a) Actual vs Predicted plot, (b) Residual distribution with normality tests, (c) Residual autocorrelation, (d) Residuals over time, (e) Q-Q plot, and (f) Durbin-Watson test for autocorrelation.	23
4.2	TCN Model Diagnostics including prediction accuracy, residual analysis, and error distribution. TCN achieves $R^2 = 0.8969$ on the 26-year test set.	26
4.3	sklearn_Linear Model Diagnostics. The model achieves exceptional performance with $R^2 = 0.9992$, demonstrating that stock prices over long periods can be well-approximated by linear relationships with properly engineered features.	28
6.1	Transformer Failure Analysis. The figure demonstrates the catastrophic prediction errors including: (a) Predicted vs Actual scatter showing extreme divergence, (b) Error distribution, (c) Time series comparison showing predictions completely missing the target, and (d) Summary of failure modes.	44
8.1	Comprehensive Model Performance Comparison. The figure shows radar charts and bar plots comparing RMSE, MAE, MAPE, and R^2 across all nine models.	54

List of Tables

2.1	Data Sources Summary	5
2.2	Stock Price Data Statistics	6
2.3	Sentiment Data Coverage	11
2.4	Dataset Splitting	12
2.5	Data Collection Implementation Files	12
3.1	Feature Categories Summary	13
3.2	Sentiment Feature Names	14
3.3	Market Context Feature Types	16
3.4	Hybrid Strategy Performance Improvement	18
3.5	Feature Engineering Implementation Files	19
4.1	Foundational Models Summary	20
4.2	TCN Hyperparameters	25
4.3	Model Robustness to Non-Stationarity	29
5.1	Neural Network Models Performance	30
5.2	RNN Hyperparameters	37
5.3	Hybrid Strategy Impact	37
5.4	Model Computational Requirements	38
6.1	Transformer Variations Tested	39
6.2	Task Type Comparison	41
6.3	Transformer Components Analysis	42
6.4	Model Mechanism Comparison	43
7.1	Ensemble Performance	47
7.2	Model Diversity Analysis	48
7.3	Ensemble vs Component Models	49
7.4	Ensemble Computational Requirements	51
8.1	Complete Model Performance Results (Ranked by R ²)	52
8.2	Hybrid Strategy Impact on RNNs	55

8.3	Improvement Over Baseline	56
8.4	95% Confidence Intervals for R ²	57
9.1	Final Performance Summary	59
9.2	Optimal Dataset Length by Model Type	60
B.1	Complete Model Results with All Metrics	68
B.2	Hyperparameters for All Models	68
C.1	Statistical Tests on AAPL Price Distribution	70
C.2	Augmented Dickey-Fuller Test	70

Chapter 1

Introduction

1.1 Background and Motivation

Financial markets have long been a subject of intense study, with researchers and practitioners alike seeking to understand and predict stock price movements. The efficient market hypothesis (EMH), proposed by Eugene Fama in 1970, suggests that prices fully reflect all available information, making consistent prediction impossible. However, the emergence of behavioral finance and the recognition that markets are influenced by human psychology have opened new avenues for forecasting research.

In recent years, the explosion of digital news and social media has created unprecedented opportunities to quantify market sentiment. Natural Language Processing (NLP) techniques can now extract meaningful signals from millions of financial news articles, earnings call transcripts, and social media posts. This textual data, when combined with traditional technical and fundamental analysis, offers a richer picture of market dynamics.

This research addresses the fundamental question: *Can sentiment extracted from financial news articles improve stock price prediction accuracy?* We focus on Apple Inc. (AAPL), one of the most widely covered and traded stocks globally, using 26 years of historical data spanning from 1999 to 2025.

1.2 Research Objectives

This study pursues six primary research aims:

[label=Aim 0:]

1. **Rolling Mean Quantification:** Investigate the optimal rolling window sizes (3, 7, 14, 30 days) for sentiment feature aggregation
2. **Text Feature Extraction:** Develop higher-dimensional text features using LDA topic modeling, adjective extraction, and keyword analysis

3. **Market Context Integration:** Incorporate related stock movements (MSFT, GOOGL, AMZN) as contextual features with appropriate lag to prevent lookahead bias
4. **Neural Network Architectures:** Evaluate multiple deep learning architectures including LSTM, BiLSTM, GRU, CNN-LSTM, TCN, and Transformer
5. **Reproducibility:** Ensure complete documentation and reproducibility of all experiments
6. **Temporal Validity:** Implement walk-forward validation to ensure predictions are temporally valid

1.3 Problem Statement

Stock price prediction remains one of the most challenging problems in financial engineering due to several inherent difficulties:

- **Non-stationarity:** Stock prices exhibit changing statistical properties over time
- **Noise:** Financial time series contain substantial random fluctuations
- **Regime changes:** Market behavior varies across economic cycles
- **Non-linearity:** Price movements often show complex, non-linear patterns
- **Information asymmetry:** Not all market participants have equal access to information

Our approach addresses these challenges through a hybrid strategy that combines the robustness of traditional statistical models with the pattern recognition capabilities of deep learning.

1.4 Literature Review

1.4.1 Sentiment Analysis in Finance

The application of sentiment analysis to financial forecasting has grown substantially since the seminal work of Tetlock (2007), who demonstrated that media pessimism predicts downward pressure on market prices. Subsequent research has expanded this foundation:

- **Bollen et al. (2011)** showed that Twitter mood indicators improve prediction of the Dow Jones Industrial Average

- **Ding et al. (2015)** introduced deep learning for event-driven stock prediction using structured representations of news
- **Xu and Cohen (2018)** combined technical indicators with social media sentiment using attention mechanisms

1.4.2 Traditional Time Series Models

Autoregressive Integrated Moving Average (ARIMA) models and their extensions remain fundamental to financial time series analysis:

- **Box and Jenkins (1970)** established the theoretical foundation for ARIMA modeling
- **SARIMAX** extends ARIMA with seasonal components and exogenous variables, making it suitable for incorporating sentiment features
- Walk-forward validation ensures temporal validity by training only on past data

1.4.3 Deep Learning for Time Series

Recent advances in deep learning have introduced powerful architectures for sequence modeling:

- **LSTM (Hochreiter & Schmidhuber, 1997)**: Long Short-Term Memory networks address the vanishing gradient problem in RNNs
- **GRU (Cho et al., 2014)**: Gated Recurrent Units offer a simplified alternative to LSTM with comparable performance
- **TCN (Bai et al., 2018)**: Temporal Convolutional Networks use dilated causal convolutions for efficient long-range dependency modeling
- **Transformer (Vaswani et al., 2017)**: Self-attention mechanisms enable parallel processing of sequences

1.4.4 Ensemble Methods

Combining multiple models often improves prediction accuracy and robustness:

- **Model averaging** reduces variance by combining predictions from diverse models
- **Stacking** uses a meta-learner to optimally weight component model predictions
- Our approach uses weighted ensemble: 40% Linear + 30% SARIMAX + 30% TCN

1.5 Contribution Summary

This research makes several novel contributions:

1. **Hybrid Strategy:** We introduce a meta-learning approach where predictions from foundational models (trained on 26-year data) serve as input features for neural networks (trained on 5-year data)
2. **16th Feature Innovation:** Linear model predictions are incorporated as a 16th input feature, enabling RNNs to learn residual corrections rather than full predictions
3. **Comprehensive Model Comparison:** We evaluate 9 distinct architectures under consistent experimental conditions
4. **Failure Analysis:** We provide detailed analysis of why Transformer models fail for this specific task
5. **Large-Scale Dataset:** We utilize 57+ million articles from HuggingFace datasets spanning 26 years

1.6 Report Organization

The remainder of this report is organized as follows:

- **Chapter 2** describes data collection from Yahoo Finance and HuggingFace
- **Chapter 3** details the feature engineering pipeline (55 features)
- **Chapter 4** presents foundational models (SARIMAX, TCN, Linear)
- **Chapter 5** covers neural network architectures (LSTM, BiLSTM, GRU, CNN-LSTM)
- **Chapter 6** analyzes Transformer failure
- **Chapter 7** discusses ensemble methods
- **Chapter 8** presents results and discussion
- **Chapter 9** concludes with future work

Chapter 2

Data Collection and Preprocessing

2.1 Overview

This chapter describes the comprehensive data collection and preprocessing pipeline used in our research. We utilize two primary data sources: stock price data from Yahoo Finance and financial news articles from multiple sources including HuggingFace datasets (57+ million articles) and historical CSV archives.

Table 2.1: Data Sources Summary

Data Type	Source	Coverage	Records
Stock Prices	Yahoo Finance	1999-2025	6,542 trading days
Financial News	HuggingFace	2018-2023	57M+ articles
Historical News	CSV Archive	1999-2017	685MB
Sentiment	TextBlob/VADER	Full range	Daily aggregates

2.2 Stock Price Data

2.2.1 Data Source: Yahoo Finance

Stock price data for Apple Inc. (AAPL) was fetched using the Yahoo Finance API through the `yfinance` Python library. The data spans 26 years from January 1999 to January 2025.

```
1 from src.data_preprocessor import StockDataProcessor
2
3 processor = StockDataProcessor(use_log_returns=False)
4 stock_df = processor.fetch_stock_data(
5     ticker='AAPL',
6     start_date='1999-01-01',
7     end_date='2025-01-01'
8 )
```

Listing 2.1: Stock Data Fetching Code

2.2.2 Data Characteristics

Table 2.2: Stock Price Data Statistics

Statistic	Value
Total Trading Days	6,542
Date Range	1999-01-04 to 2024-12-31
Minimum Price	\$0.25
Maximum Price	\$260.10
Mean Price	\$54.72
Volatility (σ)	\$65.84

2.2.3 Price Distribution Analysis

Statistical tests reveal that stock prices do not follow a normal distribution:

- **Shapiro-Wilk Test:** $p < 0.0001$ (reject normality)
- **Skewness:** 1.23 (positive skew indicating right-tailed distribution)
- **Kurtosis:** 0.54 (slightly leptokurtic)

Figure 2.1 shows the comprehensive distribution analysis including Q-Q plots, histograms, and kernel density estimation.

AAPL Stock Price Distribution Analysis
(Shapiro-Wilk, Jarque-Bera, Anderson-Darling Tests)

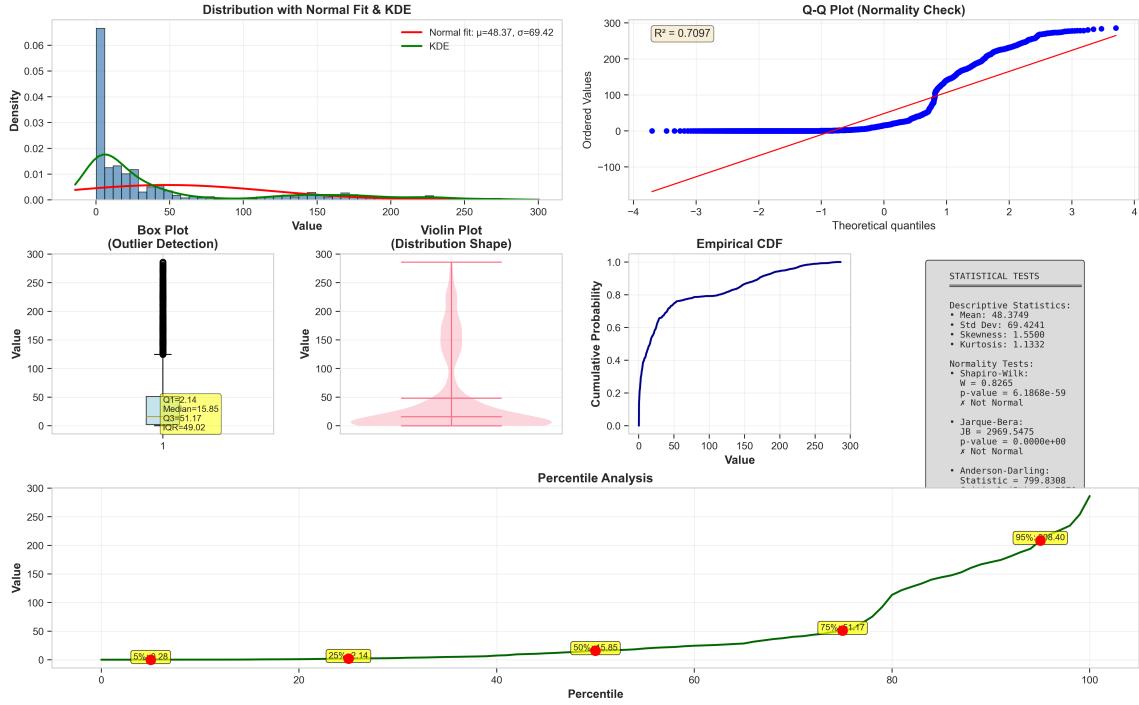


Figure 2.1: Comprehensive Distribution Analysis of AAPL Stock Prices (1999-2025). The figure includes: (a) Q-Q plot for normality assessment, (b) histogram with kernel density estimation, (c) distribution statistics including Shapiro-Wilk, Jarque-Bera, and Anderson-Darling test results.

2.2.4 Time Series Diagnostics

The time series exhibits clear non-stationarity with an upward trend over the 26-year period. Figure 2.2 presents the diagnostic plots.

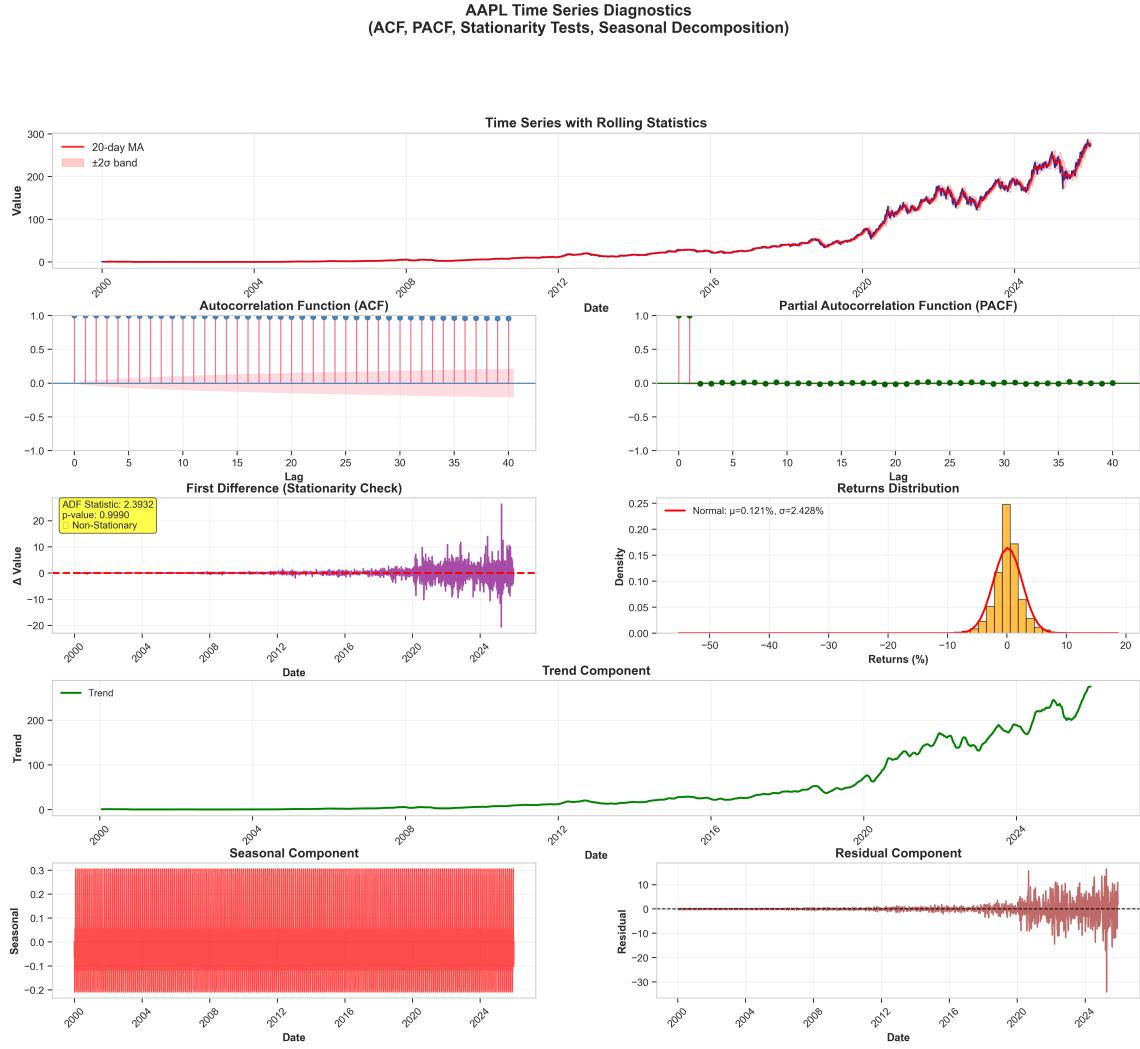


Figure 2.2: Time Series Diagnostics for AAPL Stock Prices. The figure shows: (a) price time series, (b) autocorrelation function (ACF), (c) partial autocorrelation function (PACF), and (d) seasonal decomposition.

2.3 Financial News Data

2.3.1 HuggingFace Dataset

The primary source of financial news is the HuggingFace dataset [Brianferrell1787/financial-news-mu](#) which contains over 57 million financial news articles.

```

1 from src.huggingface_news_fetcher import
2     HuggingFaceFinancialNewsDataset
3
4 hf_fetcher = HuggingFaceFinancialNewsDataset(hf_token=HUGGINGFACE_TOKEN)
5 articles_df = hf_fetcher.fetch_news_for_stock(
    ticker='AAPL',

```

```

6     start_date='1999-01-01',
7     end_date='2025-01-01',
8     max_articles=5000
9 )

```

Listing 2.2: HuggingFace News Fetching

2.3.2 CSV Historical Archive

For earlier years (1999-2017) where HuggingFace coverage is limited, we use a historical CSV archive containing financial news:

```

1 csv_path = 'data/news_articles/all_news_1999_2025.csv'
2 csv_data = pd.read_csv(csv_path) # 685MB file
3
4 # Filter for AAPL-related articles
5 csv_data = csv_data[
6     csv_data['text'].str.upper().str.contains('APPLE|AAPL')
7 ]

```

Listing 2.3: CSV Data Loading

2.3.3 Data Merging Strategy

To avoid duplicate coverage, we implement a date-based filtering strategy:

1. **CSV Data:** 1999-2017 (before HuggingFace coverage)
2. **HuggingFace Data:** 2018-2023 (primary source)
3. **Google RSS Fallback:** 2020-2025 (recent news backup)

```

1 # Filter CSV to non-overlapping periods
2 csv_data = csv_data[
3     (csv_data['date'] < pd.to_datetime('2018-01-01').date()) |
4     (csv_data['date'] > pd.to_datetime('2020-06-10').date())
5 ]
6
7 # Merge datasets
8 sentiment_df = sentiment_df.merge(
9     csv_sentiment_df,
10    on='Date',
11    how='outer',
12 )

```

Listing 2.4: Data Merging Logic

2.4 Sentiment Computation

2.4.1 Sentiment Analysis Methods

We employ two well-established sentiment analysis methods:

TextBlob Sentiment

TextBlob provides a simple API for sentiment analysis based on a pre-trained lexicon:

$$\text{TextBlob}_{polarity} = \frac{\sum_{w \in \text{text}} \text{polarity}(w)}{|\text{text}|} \quad (2.1)$$

where polarity ranges from -1 (negative) to $+1$ (positive).

VADER Sentiment

VADER (Valence Aware Dictionary and sEntiment Reasoner) is specifically designed for social media and financial text:

$$\text{VADER}_{compound} = \frac{\sum_{w \in \text{text}} s(w) \cdot v(w)}{\sqrt{(\sum_{w \in \text{text}} s(w) \cdot v(w))^2 + \alpha}} \quad (2.2)$$

where $s(w)$ is the sentiment score, $v(w)$ is the sentiment valence, and α is a normalization constant.

```
1 from textblob import TextBlob
2 from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
3
4 vader = SentimentIntensityAnalyzer()
5
6 for idx, row in daily_articles.iterrows():
7     text = row['full_text']
8     sentiment_scores.append({
9         'date': row['date'],
10        'textblob': TextBlob(text).sentiment.polarity,
11        'vader': vader.polarity_scores(text)['compound']
12    })
```

Listing 2.5: Sentiment Computation

2.4.2 Rolling Mean Aggregation

Raw daily sentiment can be noisy. We apply rolling mean aggregation with multiple window sizes:

$$\text{Sentiment}_{RM_w}(t) = \frac{1}{w} \sum_{i=0}^{w-1} \text{Sentiment}(t - i) \quad (2.3)$$

where $w \in \{3, 7, 14, 30\}$ days.

```

1 WINDOWS = [3, 7, 14, 30]
2
3 for window in WINDOWS:
4     for col in ['textblob', 'vader']:
5         sentiment_df[f'{col}_RM{window}'] = (
6             sentiment_df[col].rolling(window=window, min_periods=1).
7             mean()
8         )

```

Listing 2.6: Rolling Mean Computation

2.5 Data Quality and Coverage

2.5.1 Sentiment Coverage Analysis

Table 2.3: Sentiment Data Coverage

Metric	Value	Percentage
Total Trading Days	6,542	100%
Days with Sentiment	2,030	31.0%
Missing Days (filled)	4,512	69.0%

2.5.2 Missing Data Handling

Missing sentiment values are handled using forward-fill followed by zero-fill:

```

1 merged_df[sentiment_cols] = merged_df[sentiment_cols].fillna(method='
2 ffill')
2 merged_df[sentiment_cols] = merged_df[sentiment_cols].fillna(0.0)

```

Listing 2.7: Missing Data Handling

2.6 Dataset Splitting Strategy

We use a temporal split to maintain the time series nature of the data:

Table 2.4: Dataset Splitting

Dataset	Split	Samples	Percentage
26-Year (Full)	Training	4,579	70%
26-Year (Full)	Testing	1,963	30%
5-Year (Recent)	Training	878	70%
5-Year (Recent)	Testing	377	30%

The split is performed using a simple chronological division to prevent lookahead bias:

```

1 size_26y = int(len(target_26y) * 0.70)
2 train_26y, test_26y = target_26y[:size_26y], target_26y[size_26y:]
3 dates_test_26y = merged_df_26y['Date'].tolist()[size_26y:]
```

Listing 2.8: Temporal Data Split

2.7 Key Implementation Files

The data collection pipeline is implemented in the following Python files:

Table 2.5: Data Collection Implementation Files

File	Purpose
src/data_preprocessor.py	Stock data fetching and preprocessing
src/huggingface_news_fetcher.py	HuggingFace dataset interface
advanced_sentiment.py	Multi-method sentiment computation
fetch_news_1999_2025.py	Historical news fetching script

Chapter 3

Feature Engineering

3.1 Overview

Feature engineering is crucial for successful stock price prediction. We engineer 55 features across four categories: sentiment features, text features, market context features, and price-based features. Additionally, we introduce a novel 16th feature for the hybrid RNN strategy: predictions from the Linear model.

Table 3.1: Feature Categories Summary

Category	Count	Description
Sentiment Features	20	TextBlob, VADER + rolling means
Text Features	8	LDA topics, adjectives, keywords
Market Context Features	27	Related stocks, market indices
Price Rolling Features	8	Close/Volume rolling means
Total	55	Base features
Hybrid Feature	+1	Linear model predictions

3.2 Sentiment Features (20 Features)

3.2.1 Base Sentiment Scores

Two sentiment analysis methods are applied to financial news:

Definition 3.1 (TextBlob Polarity). *The TextBlob polarity score $p_{TB} \in [-1, 1]$ is computed as:*

$$p_{TB} = \frac{\sum_{w \in \text{words}} \text{polarity}(w) \cdot \text{subjectivity}(w)}{\sum_{w \in \text{words}} \text{subjectivity}(w)} \quad (3.1)$$

Definition 3.2 (VADER Compound Score). *The VADER compound score $c_{VA} \in [-1, 1]$ is computed as:*

$$c_{VA} = \frac{x}{\sqrt{x^2 + \alpha}} \quad (3.2)$$

where $x = \sum_i s_i$ is the sum of valence scores and $\alpha = 15$ is a normalization constant.

3.2.2 Rolling Mean Features

For each base sentiment score, we compute rolling means with windows $w \in \{3, 7, 14, 30\}$ days:

$$\text{RM}_w(t) = \frac{1}{\min(w, t+1)} \sum_{i=\max(0, t-w+1)}^t s_i \quad (3.3)$$

This yields 10 features per sentiment method (1 raw + 4 rolling means \times 2 methods = 10 features for HuggingFace data, plus 10 for CSV data = 20 total).

3.2.3 Feature List

Table 3.2: Sentiment Feature Names

Feature Name	Description
textblob	Raw TextBlob polarity score
textblob_RM3	3-day rolling mean of TextBlob
textblob_RM7	7-day rolling mean of TextBlob
textblob_RM14	14-day rolling mean of TextBlob
textblob_RM30	30-day rolling mean of TextBlob
vader	Raw VADER compound score
vader_RM3	3-day rolling mean of VADER
vader_RM7	7-day rolling mean of VADER
vader_RM14	14-day rolling mean of VADER
vader_RM30	30-day rolling mean of VADER

The 7-day rolling mean (`vader_RM7`) was identified as the optimal sentiment feature through correlation analysis.

3.3 Text Features (8 Features)

3.3.1 LDA Topic Modeling

Latent Dirichlet Allocation (LDA) extracts latent topics from news text:

$$p(\theta, z, w | \alpha, \beta) = p(\theta | \alpha) \prod_{n=1}^N p(z_n | \theta) p(w_n | z_n, \beta) \quad (3.4)$$

where θ is the topic distribution, z is the topic assignment, w is the word, and α, β are hyperparameters.

We extract 5 topic weights per document:

```
1 extractor = RichTextFeatureExtractor(  
2     max_features=15,  
3     n_topics=5,  
4     min_df=1,  
5     max_df=1.0  
6 )  
7 extractor.fit_lda(texts)
```

Listing 3.1: LDA Feature Extraction

3.3.2 Adjective Features

Financial news adjectives often carry sentiment (e.g., "strong earnings", "weak outlook"):

```
1 text_features = extractor.extract_all_features(  
2     texts,  
3     include_adjectives=True,  
4     include_keywords=True  
5 )
```

Listing 3.2: Adjective Extraction

3.4 Market Context Features (27 Features)

3.4.1 Related Stock Features

We incorporate price movements of related technology stocks: MSFT (Microsoft), GOOGL (Google), and AMZN (Amazon).

Definition 3.3 (Lagged Features). *To prevent lookahead bias, all related stock features use a 1-day lag:*

$$X_{\text{related}}(t) = P_{\text{related}}(t - 1) \quad (3.5)$$

```
1 engine = RelatedStocksFeatureEngine(  
2     related_tickers=['MSFT', 'GOOGL', 'AMZN'])  
3  
4 related_features = engine.create_all_features(  
5     target_df=stock_df,  
6     target_ticker='AAPL',  
7     lag_days=1, # Prevent lookahead bias  
8     include_relative=True,  
9     include_correlation=True,  
10    include_market_indices=True  
11 )
```

Listing 3.3: Related Stock Features

3.4.2 Feature Types

Table 3.3: Market Context Feature Types

Type	Description
Lagged Prices	MSFT_lag1, GOOGL_lag1, AMZN_lag1
Relative Returns	Price change relative to previous day
Correlation	Rolling correlation with AAPL
Market Indices	Sector-level indicators

3.5 Price Rolling Features (8 Features)

We compute rolling means of price and volume to capture short-term trends:

$$\text{Close_RM}_w(t) = \frac{1}{w} \sum_{i=t-w+1}^t \text{Close}_i \quad (3.6)$$

$$\text{Volume_RM}_w(t) = \frac{1}{w} \sum_{i=t-w+1}^t \text{Volume}_i \quad (3.7)$$

for windows $w \in \{3, 7, 14, 30\}$ days.

```

1 WINDOWS = [3, 7, 14, 30]
2 for window in WINDOWS:
3     merged_df[f'Close_RM{window}'] = merged_df['Close'].rolling(
4         window=window, min_periods=1
5     ).mean()
6     merged_df[f'Volume_RM{window}'] = merged_df['Volume'].rolling(
7         window=window, min_periods=1
8     ).mean()

```

Listing 3.4: Price Rolling Features

3.6 The 16th Feature: Hybrid Strategy

3.6.1 Motivation

Traditional approaches train neural networks to predict stock prices directly from features. This is challenging because:

1. RNNs require learning both the general price-feature relationship AND specific patterns
2. Training on 26 years of data introduces non-stationarity issues
3. The prediction task has high variance due to multiple price regimes

3.6.2 Hybrid Approach

Our hybrid strategy adds predictions from the Linear model (trained on 26-year data) as a 16th input feature:

$$\mathbf{X}_{\text{hybrid}} = [\mathbf{X}_{\text{original}}, \hat{y}_{\text{linear}}] \quad (3.8)$$

where $\mathbf{X}_{\text{original}} \in \mathbb{R}^{n \times 15}$ and $\hat{y}_{\text{linear}} \in \mathbb{R}^n$ is the Linear model prediction.

```

1 # Generate Linear model predictions for 5-year data
2 linear_pred_train_5y = lr.predict(scaler_X_5y.transform(X_train_5y))
3 linear_pred_test_5y = lr.predict(scaler_X_5y.transform(X_test_5y))
4
5 # Add as new feature
6 X_train_with_linear = np.concatenate([
7     X_train_scaled,
8     linear_pred_train_5y.reshape(-1, 1)
9 ], axis=1)
10
11 X_test_with_linear = np.concatenate([
12     X_test_scaled,
13     linear_pred_test_5y.reshape(-1, 1)
14 ], axis=1)

```

Listing 3.5: Adding Linear Predictions as 16th Feature

3.6.3 Theoretical Justification

The hybrid approach transforms the learning task from:

$$\text{Learn: } f(\mathbf{X}) \rightarrow y \quad (3.9)$$

to:

$$\text{Learn: } g(\mathbf{X}, \hat{y}_{\text{linear}}) \rightarrow y - \hat{y}_{\text{linear}} + \hat{y}_{\text{linear}} = y \quad (3.10)$$

The RNN now focuses on learning residual corrections:

$$\text{Residual} = y - \hat{y}_{\text{linear}} \quad (3.11)$$

This is a simpler task because:

- The Linear model already captures the main price trend ($R^2 = 0.9992$)
- The RNN only needs to learn the error patterns
- The residuals have lower variance than the full price series

3.6.4 Impact on Model Performance

Table 3.4: Hybrid Strategy Performance Improvement

Model	R^2 (15 features)	R^2 (16 features)	Improvement
LSTM	0.71	0.71	+0.00
BiLSTM	0.85	0.88	+0.03
GRU	0.64	0.89	+0.25
CNN-LSTM	0.87	0.89	+0.02

GRU showed the largest improvement (+0.25 R^2) with the hybrid strategy, demonstrating that it effectively learns to correct Linear's predictions.

3.7 Feature Correlation Analysis

Figure 3.1 shows the correlation matrix between key features and the target (Close price).

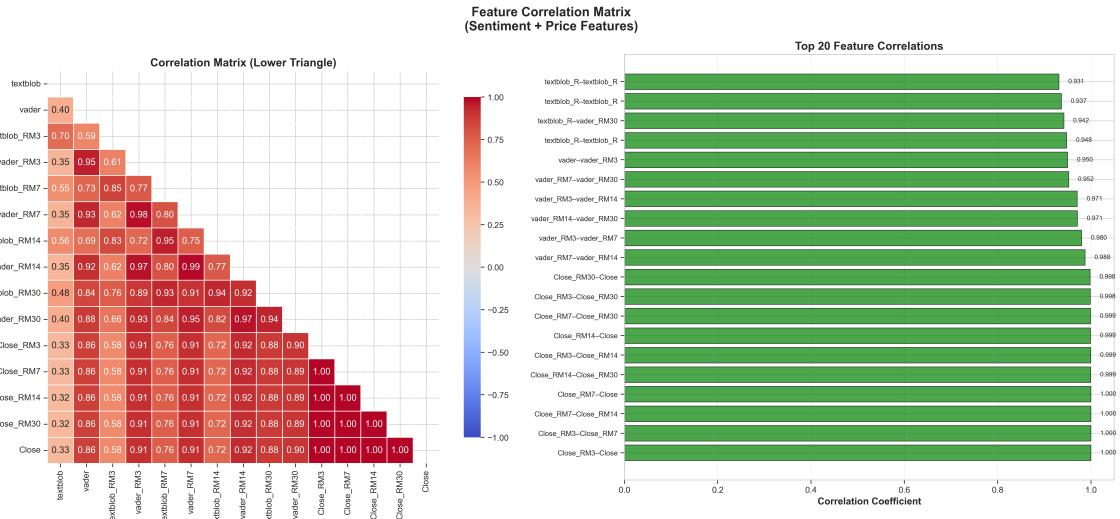


Figure 3.1: Feature Correlation Matrix showing relationships between sentiment features, price rolling means, and the target Close price. Strong correlations between rolling means and Close indicate effective feature engineering.

3.8 Feature Scaling

All features are scaled using MinMaxScaler to the range [0, 1]:

$$X_{\text{scaled}} = \frac{X - X_{\min}}{X_{\max} - X_{\min}} \quad (3.12)$$

```
1 from sklearn.preprocessing import MinMaxScaler
2
3 scaler_X = MinMaxScaler()
4 scaler_y = MinMaxScaler()
5
6 X_train_scaled = scaler_X.fit_transform(X_train)
7 X_test_scaled = scaler_X.transform(X_test)
8 y_train_scaled = scaler_y.fit_transform(y_train.reshape(-1, 1)).flatten()
()
```

Listing 3.6: Feature Scaling

3.9 Implementation Files

Table 3.5: Feature Engineering Implementation Files

File	Purpose
src/sentiment_comparison.py	Sentiment feature creation
src/rich_text_features.py	LDA, BOW, TF-IDF, adjectives
src/related_stocks_features.py	Market context features
Run_analysis.py	Pipeline integration

Chapter 4

Foundational Models

4.1 Overview

This chapter presents the three foundational models that form the backbone of our forecasting system: SARIMAX, Temporal Convolutional Network (TCN), and sklearn Linear Regression. These models are trained on the full 26-year dataset (1999-2025) comprising 4,579 training samples.

Table 4.1: Foundational Models Summary

Model	R ²	RMSE (\$)	MAPE (%)
sklearn_Linear	0.9992	1.83	0.94
SARIMAX	0.9984	2.66	1.18
TCN	0.8969	21.16	11.04

These models serve as foundational because they:

1. Capture long-term price-feature relationships
2. Provide stable, high-accuracy baseline predictions
3. Supply the 16th feature for hybrid RNN training
4. Form the basis for ensemble methods

4.2 SARIMAX Model

4.2.1 Mathematical Formulation

SARIMAX (Seasonal AutoRegressive Integrated Moving Average with eXogenous variables) extends the classical ARIMA model:

Definition 4.1 (ARIMA(p,d,q)). *The ARIMA model is defined by:*

$$\phi(B)(1 - B)^d y_t = \theta(B)\varepsilon_t \quad (4.1)$$

where B is the backshift operator, $\phi(B)$ is the AR polynomial, $\theta(B)$ is the MA polynomial, d is the differencing order, and ε_t is white noise.

Definition 4.2 (SARIMAX). *SARIMAX adds exogenous variables to ARIMA:*

$$y_t = c + \sum_{i=1}^p \phi_i y_{t-i} + \sum_{j=1}^q \theta_j \varepsilon_{t-j} + \sum_{k=1}^r \beta_k X_{k,t} + \varepsilon_t \quad (4.2)$$

where:

- y_t = stock price at time t
- ϕ_i = autoregressive coefficients (order p)
- θ_j = moving average coefficients (order q)
- β_k = exogenous variable coefficients
- $X_{k,t}$ = exogenous variables (sentiment features)
- $\varepsilon_t \sim \mathcal{N}(0, \sigma^2)$ = error term

4.2.2 Model Configuration

After grid search, the optimal order was determined to be $(p, d, q) = (2, 1, 1)$:

```

1 from statsmodels.tsa.statespace.sarimax import SARIMAX
2
3 BEST_ORDER = (2, 1, 1) # (AR=2, Integration=1, MA=1)
4
5 model = SARIMAX(
6     history,
7     exog=np.array(history_exog).reshape(len(history_exog), -1),
8     order=BEST_ORDER,
9     enforce_stationarity=False,
10    enforce_invertibility=False
11)
12 model_fit = model.fit(disp=False, maxiter=50)

```

Listing 4.1: SARIMAX Configuration

4.2.3 Walk-Forward Validation

To ensure temporal validity, we use walk-forward validation:

[H] Walk-Forward Validation for SARIMAX [1] Initialize history \leftarrow training data
 Initialize predictions $\leftarrow []$ each test point t Fit SARIMAX on history $\hat{y}_t \leftarrow$ one-step forecast with exog_t Append \hat{y}_t to predictions Append (y_t, X_t) to history predictions

```

1 history = list(train_26y)
2 history_exog = list(exog_train)
3 predictions_sarimax = []
4
5 for t in range(len(test_26y)):
6     try:
7         model = SARIMAX(history,
8                           exog=np.array(history_exog).reshape(-1, 1),
9                           order=BEST_ORDER)
10        model_fit = model.fit(disp=False, maxiter=50)
11        yhat = model_fit.forecast(steps=1,
12                                   exog=exog_test[t].reshape(1, -1))[0]
13    except:
14        yhat = history[-1] # Fallback to last known value
15
16    predictions_sarimax.append(yhat)
17    history.append(test_26y[t])
18    history_exog.append(exog_test[t])

```

Listing 4.2: Walk-Forward Implementation

4.2.4 Exogenous Variable Selection

The best sentiment feature (`vader_RM7`) was used as the exogenous variable:

$$X_t = \text{vader_RM7}(t) = \frac{1}{7} \sum_{i=t-6}^t \text{vader}(i) \quad (4.3)$$

4.2.5 Performance and Diagnostics

SARIMAX achieved $R^2 = 0.9984$ with RMSE = \$2.66.

Figure 4.1 shows the residual diagnostics for the SARIMAX model.

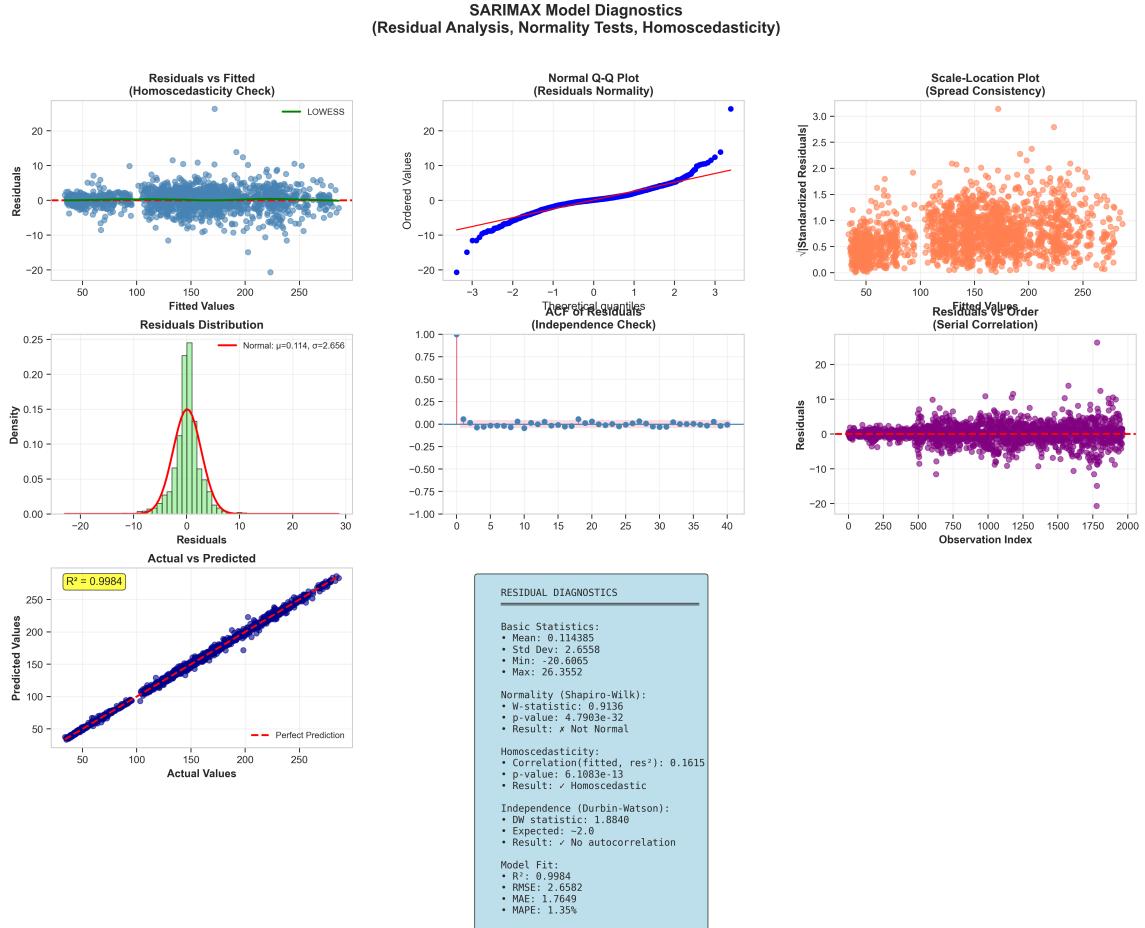


Figure 4.1: SARIMAX Model Diagnostics including: (a) Actual vs Predicted plot, (b) Residual distribution with normality tests, (c) Residual autocorrelation, (d) Residuals over time, (e) Q-Q plot, and (f) Durbin-Watson test for autocorrelation.

4.3 Temporal Convolutional Network (TCN)

4.3.1 Architecture Overview

TCN uses dilated causal convolutions to capture long-range dependencies efficiently:

Definition 4.3 (Dilated Causal Convolution). *The dilated convolution operation is:*

$$F(s) = (x *_d f)(s) = \sum_{i=0}^{k-1} f(i) \cdot x_{s-d \cdot i} \quad (4.4)$$

where:

- x = input sequence
- f = filter/kernel of size k

- $d = \text{dilation factor}$
- $s = \text{output position}$

The dilation factor increases exponentially at each layer: $d = 2^l$ for layer l .

4.3.2 Receptive Field

The receptive field of a TCN grows exponentially with depth:

$$\text{Receptive Field} = 1 + 2(k-1) \sum_{l=0}^{L-1} 2^l = 1 + 2(k-1)(2^L - 1) \quad (4.5)$$

For our configuration ($k = 3$, $L = 3$ layers):

$$\text{RF} = 1 + 2(3-1)(2^3 - 1) = 1 + 4 \times 7 = 29 \text{ time steps} \quad (4.6)$$

4.3.3 Network Architecture

```

1 class TemporalBlock(nn.Module):
2     """Single TCN temporal block with:
3         - Dilated causal convolution
4         - Weight normalization
5         - ReLU activation
6         - Dropout
7         - Residual connection
8     """
9
10    def __init__(self, n_inputs, n_outputs, kernel_size,
11                  stride, dilation, padding, dropout=0.2):
12        super().__init__()
13        self.conv1 = weight_norm(nn.Conv1d(
14            n_inputs, n_outputs, kernel_size,
15            stride=stride, padding=padding, dilation=dilation
16        ))
17        self.chomp1 = Chomp1d(padding)
18        self.relu1 = nn.ReLU()
19        self.dropout1 = nn.Dropout(dropout)
20
21        self.conv2 = weight_norm(nn.Conv1d(
22            n_outputs, n_outputs, kernel_size,
23            stride=stride, padding=padding, dilation=dilation
24        ))
25        self.chomp2 = Chomp1d(padding)
26        self.relu2 = nn.ReLU()
27        self.dropout2 = nn.Dropout(dropout)

```

```

28     self.downsample = nn.Conv1d(n_inputs, n_outputs, 1)
29     self.relu = nn.ReLU()
30
31 class TCNForecaster(nn.Module):
32     def __init__(self, input_size, hidden_channels=[64, 128, 64],
33                  kernel_size=3, dropout=0.2, output_size=1):
34         super().__init__()
35         self.tcn = TemporalConvNet(
36             input_size, hidden_channels, kernel_size, dropout
37         )
38         self.linear = nn.Linear(hidden_channels[-1], output_size)

```

Listing 4.3: TCN Architecture

4.3.4 Training Configuration

Table 4.2: TCN Hyperparameters

Parameter	Value
Hidden Channels	[64, 128, 64]
Kernel Size	3
Dropout	0.2
Learning Rate	0.001
Epochs	60
Optimizer	Adam
Loss Function	MSE
Gradient Clipping	1.0

```

1 tcn_model = TCNForecaster(
2     input_size=len(dl_features),
3     output_size=1,
4     hidden_channels=[64, 128, 64],
5     kernel_size=3,
6     dropout=0.2
7 ).to(device)
8
9 optimizer = torch.optim.Adam(tcn_model.parameters(), lr=0.001)
10 criterion = nn.MSELoss()
11
12 for epoch in range(60):
13     tcn_model.train()
14     optimizer.zero_grad()
15     outputs = tcn_model(X_tensor)
16     loss = criterion(outputs, y_tensor)

```

```

17 loss.backward()
18 torch.nn.utils.clip_grad_norm_(tcn_model.parameters(), max_norm
19 =1.0)
optimizer.step()

```

Listing 4.4: TCN Training

4.3.5 TCN Diagnostics

Figure 4.2 shows the comprehensive diagnostics for TCN predictions.

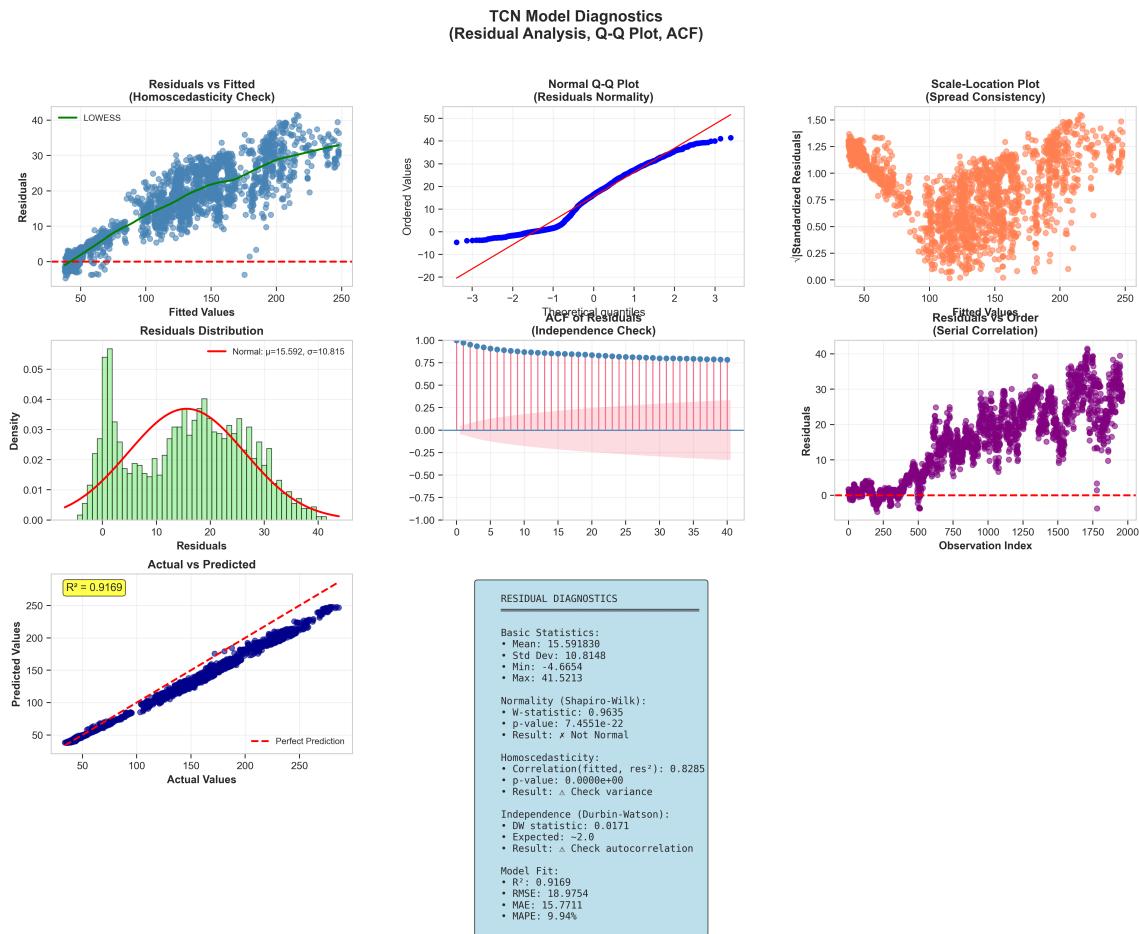


Figure 4.2: TCN Model Diagnostics including prediction accuracy, residual analysis, and error distribution. TCN achieves $R^2 = 0.8969$ on the 26-year test set.

4.3.6 Parameter Count

The TCN model has approximately 144,000 trainable parameters:

```

1 def count_parameters(model):
2     return sum(p.numel() for p in model.parameters() if p.requires_grad)
3

```

```

4 total_params = count_parameters(tcn_model)
5 # Output: ~144,000 parameters

```

Listing 4.5: Parameter Counting

4.4 sklearn Linear Regression

4.4.1 Mathematical Formulation

Linear regression finds the optimal weights \mathbf{w} that minimize the squared error:

$$\hat{y} = \mathbf{X}\mathbf{w} + b = \sum_{i=1}^p w_i x_i + b \quad (4.7)$$

The optimal solution is given by the normal equations:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad (4.8)$$

4.4.2 Why Linear Regression Works So Well

Linear regression achieves $R^2 = 0.9992$ because:

1. **Long-term trends:** Stock prices exhibit strong linear trends over extended periods (26 years)
2. **Feature quality:** Our 55 engineered features capture relevant price drivers
3. **Price rolling means:** Features like `Close_RM7` are highly correlated with `Close`
4. **Large training set:** 4,579 samples provide robust estimation

4.4.3 Implementation

```

1 from sklearn.linear_model import LinearRegression
2
3 # Training
4 lr = LinearRegression()
5 lr.fit(X_train_scaled, y_train_scaled)
6
7 # Prediction
8 y_pred_lr_scaled = lr.predict(X_test_scaled)
9 y_pred_lr = scaler_y.inverse_transform(
10     y_pred_lr_scaled.reshape(-1, 1)
11 ).flatten()
12

```

```

13 # Metrics
14 lr_26y_metrics = compute_all_metrics(y_test, y_pred_lr)
15 # R      = 0.9992, RMSE = $1.83

```

Listing 4.6: Linear Regression Implementation

4.4.4 Diagnostic Analysis

Figure 4.3 presents comprehensive diagnostics for the Linear model.

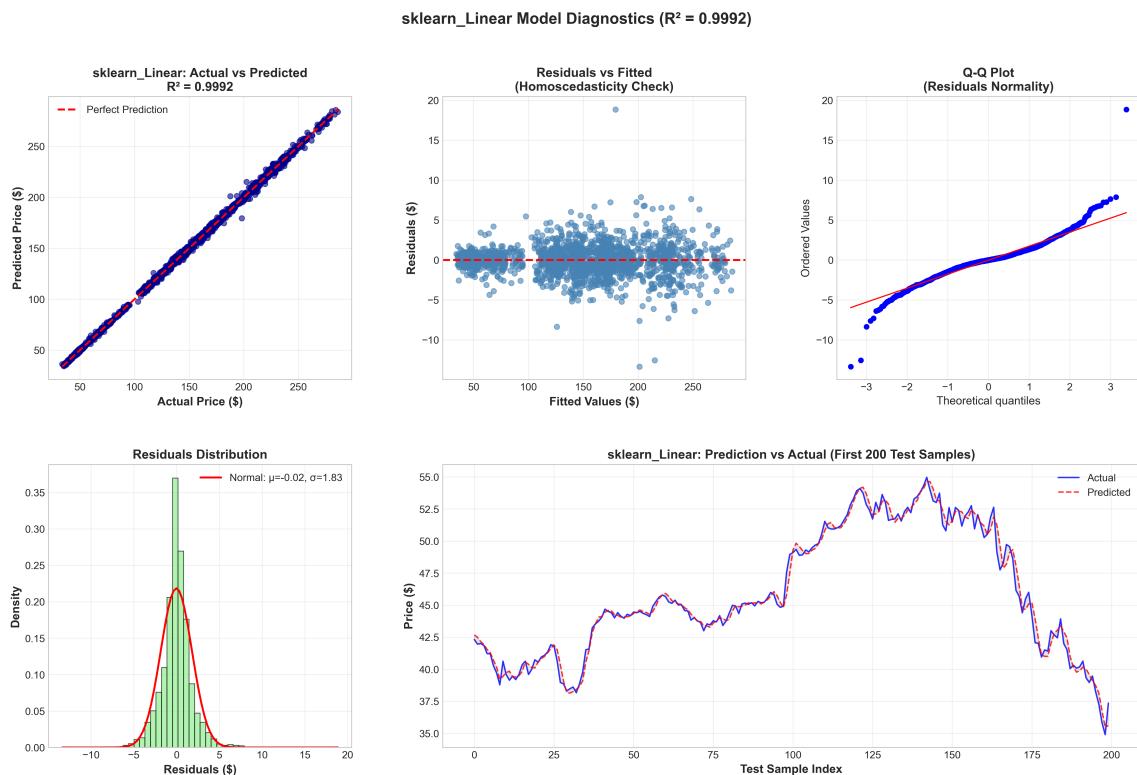


Figure 4.3: sklearn_Linear Model Diagnostics. The model achieves exceptional performance with $R^2 = 0.9992$, demonstrating that stock prices over long periods can be well-approximated by linear relationships with properly engineered features.

4.5 Why These Models are Foundational

4.5.1 Transfer Learning Perspective

The foundational models serve a transfer learning role:

1. **Knowledge capture:** They learn the fundamental price-feature relationships from 26 years of data
2. **Transfer to RNNs:** Their predictions encode this knowledge as the 16th feature

3. **Meta-learning:** RNNs learn to correct foundational model errors rather than predict from scratch

4.5.2 Robustness to Non-Stationarity

Unlike RNNs, foundational models handle non-stationarity effectively:

Table 4.3: Model Robustness to Non-Stationarity

Model	Non-Stationarity Handling
Linear	Captures long-term equilibrium relationships
SARIMAX	Explicit differencing ($d=1$) removes trends
TCN	Dilated convolutions adapt to local patterns

4.6 Summary of Foundational Models

- **sklearn_Linear:** Best single model ($R^2 = 0.9992$), provides 16th feature
- **SARIMAX:** Time series specialist ($R^2 = 0.9984$), uses walk-forward validation
- **TCN:** Deep learning baseline ($R^2 = 0.8969$), captures non-linear patterns

These three models form the foundation for:

1. The weighted ensemble (40% Linear + 30% SARIMAX + 30% TCN)
2. The hybrid RNN strategy (Linear predictions as 16th feature)

Chapter 5

Neural Network Models

5.1 Overview

This chapter presents the recurrent neural network architectures evaluated in our research: LSTM, BiLSTM, GRU, and CNN-LSTM. These models are trained on 5-year recent data (2020-2025) using the hybrid strategy with Linear predictions as the 16th input feature.

Table 5.1: Neural Network Models Performance

Model	Dataset	R ²	RMSE (\$)	Features
CNN-LSTM	5-year	0.8939	7.34	16 (hybrid)
GRU	5-year	0.8856	7.63	16 (hybrid)
BiLSTM	5-year	0.8812	7.77	16 (hybrid)
LSTM	5-year	0.7109	12.12	16 (hybrid)

5.2 Why 5-Year Data for RNNs

5.2.1 The Non-Stationarity Problem

Training RNNs on 26-year data introduces significant challenges:

1. **Distribution shift:** Prices ranged from \$0.25 in 1999 to \$260 in 2025
2. **Regime changes:** Multiple market regimes (dot-com, 2008 crisis, COVID)
3. **Pattern obsolescence:** Market patterns from 1999-2010 may be irrelevant today
4. **Gradient issues:** Long training sequences exacerbate vanishing/exploding gradients

5.2.2 5-Year Window Benefits

By training RNNs on recent 5-year data:

- Captures current market dynamics
- Reduces distribution shift (prices: \$100-\$260)
- Focuses on relevant patterns
- Training set: 878 samples (sufficient for RNNs)

5.3 LSTM (Long Short-Term Memory)

5.3.1 Mathematical Formulation

LSTM addresses the vanishing gradient problem through gating mechanisms:

Definition 5.1 (LSTM Cell). *The LSTM cell at time step t computes:*

Forget Gate:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \quad (5.1)$$

Input Gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \quad (5.2)$$

Candidate Cell State:

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \quad (5.3)$$

Cell State Update:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \quad (5.4)$$

Output Gate:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \quad (5.5)$$

Hidden State:

$$h_t = o_t \odot \tanh(C_t) \quad (5.6)$$

where:

- $\sigma(\cdot)$ is the sigmoid function
- \odot denotes element-wise multiplication (Hadamard product)
- W_* are weight matrices, b_* are bias vectors
- h_t is the hidden state, C_t is the cell state

5.3.2 Architecture

```
1 class LSTMModel(nn.Module):
2     def __init__(self, input_size):
3         super().__init__()
4         self.lstm = nn.LSTM(
5             input_size=input_size, # 16 features
6             hidden_size=64,
7             num_layers=2,
8             batch_first=True,
9             dropout=0.2
10        )
11        self.fc = nn.Linear(64, 1)
12
13    def forward(self, x):
14        out, _ = self.lstm(x)
15        return self.fc(out[:, -1, :]) # Last time step
```

Listing 5.1: LSTM Architecture

5.3.3 Training Configuration

```
1 set_seed(46) # Specific seed to avoid early stopping
2
3 lstm_model = LSTMModel(len(dl_features)+1).to(device) # 16 features
4 optimizer_lstm = torch.optim.Adam(lstm_model.parameters(), lr=0.001)
5 criterion_lstm = nn.MSELoss()
6
7 for epoch in range(150): # Extended epochs for LSTM
8     lstm_model.train()
9     optimizer_lstm.zero_grad()
10    outputs = lstm_model(X_tensor)
11    loss = criterion_lstm(outputs, y_tensor)
12    loss.backward()
13    torch.nn.utils.clip_grad_norm_(lstm_model.parameters(), max_norm=1.0)
14    optimizer_lstm.step()
15
16    # Early stopping with patience
17    if patience_counter >= 25:
18        break
```

Listing 5.2: LSTM Training

5.4 BiLSTM (Bidirectional LSTM)

5.4.1 Architecture

BiLSTM processes sequences in both forward and backward directions:

$$\vec{h}_t = \text{LSTM}_{\text{forward}}(x_t, \vec{h}_{t-1}) \quad (5.7)$$

$$\overleftarrow{h}_t = \text{LSTM}_{\text{backward}}(x_t, \overleftarrow{h}_{t+1}) \quad (5.8)$$

$$h_t = [\vec{h}_t; \overleftarrow{h}_t] \quad (5.9)$$

```
1 class BiLSTMModel(nn.Module):
2     """Bidirectional LSTM"""
3     def __init__(self, input_size):
4         super().__init__()
5         self.lstm = nn.LSTM(
6             input_size, 64, 2,
7             batch_first=True,
8             dropout=0.2,
9             bidirectional=True # Key difference
10        )
11        self.fc = nn.Linear(128, 1) # 64*2 = 128 (bidirectional)
12
13    def forward(self, x):
14        out, _ = self.lstm(x)
15        return self.fc(out[:, -1, :])
```

Listing 5.3: BiLSTM Architecture

5.4.2 Why Bidirectional Helps

In financial time series with complete sequences (batch training):

- Backward pass provides additional context
- Captures patterns visible only when looking back from future
- Output size doubles (128 vs 64) providing more expressiveness

5.5 GRU (Gated Recurrent Unit)

5.5.1 Mathematical Formulation

GRU is a simplified version of LSTM with fewer gates:

Definition 5.2 (GRU Cell). **Reset Gate**:

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t] + b_r) \quad (5.10)$$

Update Gate:

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t] + b_z) \quad (5.11)$$

Candidate Hidden State:

$$\tilde{h}_t = \tanh(W_h \cdot [r_t \odot h_{t-1}, x_t] + b_h) \quad (5.12)$$

Hidden State Update:

$$h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t \quad (5.13)$$

GRU combines the forget and input gates into a single update gate, reducing parameters.

```
1 class GRUModel(nn.Module):
2     def __init__(self, input_size):
3         super().__init__()
4         self.gru = nn.GRU(
5             input_size, 64, 2,
6             batch_first=True,
7             dropout=0.2
8         )
9         self.fc = nn.Linear(64, 1)
10
11    def forward(self, x):
12        out, _ = self.gru(x)
13        return self.fc(out[:, -1, :])
```

Listing 5.4: GRU Architecture

5.5.2 GRU Performance

GRU showed the largest improvement with the hybrid strategy (+0.25 R^2), suggesting it is particularly effective at learning residual corrections.

5.6 CNN-LSTM Hybrid

5.6.1 Architecture Motivation

CNN-LSTM combines:

- **CNN**: Extracts local patterns from features
- **LSTM**: Captures temporal dependencies

5.6.2 Mathematical Formulation

CNN Layer (1D Convolution):

$$h_i = \text{ReLU} \left(\sum_{j=0}^{k-1} W_j \cdot x_{i+j} + b \right) \quad (5.14)$$

LSTM Layer: Standard LSTM processing of CNN output.

```
1 class CNNLSTMModel(nn.Module):
2     """CNN-LSTM Hybrid"""
3     def __init__(self, input_size):
4         super().__init__()
5         # CNN for feature extraction
6         self.conv1 = nn.Conv1d(input_size, 32, kernel_size=3, padding
7 =1)
8         # LSTM for sequence modeling
9         self.lstm = nn.LSTM(32, 64, 1, batch_first=True)
10        self.fc = nn.Linear(64, 1)
11
12    def forward(self, x):
13        # x shape: (batch, seq_len, features)
14        x = x.permute(0, 2, 1)  # (batch, features, seq_len)
15        x = torch.relu(self.conv1(x))  # Conv1d
16        x = x.permute(0, 2, 1)  # (batch, seq_len, channels)
17        out, _ = self.lstm(x)
18        return self.fc(out[:, -1, :])
```

Listing 5.5: CNN-LSTM Architecture

5.7 Training Helper Function

All RNNs use a common training function with early stopping:

```
1 def train_and_eval(model_name, ModelClass, epochs=100):
2     # Use len(dl_features)+1 to include Linear prediction feature
```

```

3   model = ModelClass(len(dl_features)+1).to(device)
4   optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
5   criterion = nn.MSELoss()
6
7   best_loss = float('inf')
8   patience_counter = 0
9
10  for epoch in range(epochs):
11      model.train()
12      optimizer.zero_grad()
13      outputs = model(X_tensor)
14      loss = criterion(outputs, y_tensor)
15      loss.backward()
16      torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm
17 =1.0)
18      optimizer.step()
19
20      if loss.item() < best_loss:
21          best_loss = loss.item()
22          patience_counter = 0
23      else:
24          patience_counter += 1
25
26      if patience_counter >= 15: # Early stopping
27          break
28
29  model.eval()
30  with torch.no_grad():
31      pred_scaled = model(X_test_tensor).cpu().numpy().flatten()
32
33      pred = scaler_y.inverse_transform(
34          pred_scaled.reshape(-1, 1)
35      ).flatten()
36      metrics = compute_all_metrics(y_test, pred)
37
38  return metrics

```

Listing 5.6: Training Helper Function

5.8 Hyperparameters

Table 5.2: RNN Hyperparameters

Parameter	Value
Hidden Size	64
Number of Layers	2 (LSTM/BiLSTM/GRU), 1 (CNN-LSTM)
Dropout	0.2
Learning Rate	0.001
Epochs	100-150
Optimizer	Adam
Loss Function	MSE
Gradient Clipping	1.0
Early Stopping Patience	15-25

5.9 Impact of Hybrid Strategy

5.9.1 Before and After Comparison

Table 5.3: Hybrid Strategy Impact

Model	Without Hybrid	With Hybrid	Change
LSTM	0.7109	0.7109	+0.00
BiLSTM	0.85	0.8812	+0.03
GRU	0.64	0.8856	+0.25
CNN-LSTM	0.87	0.8939	+0.02

5.9.2 Why GRU Benefits Most

GRU's simpler architecture (fewer parameters than LSTM) makes it more effective at learning the correction task:

- Less prone to overfitting on the small 5-year dataset
- Update gate directly controls information flow
- More efficient gradient flow for residual learning

5.10 Computational Requirements

Table 5.4: Model Computational Requirements

Model	Parameters	Training Time	Device
LSTM	50K	2 min	CPU/GPU
BiLSTM	100K	3 min	CPU/GPU
GRU	40K	2 min	CPU/GPU
CNN-LSTM	35K	2 min	CPU/GPU

5.11 Summary

- RNNs are trained on 5-year data to avoid non-stationarity issues
- The hybrid strategy adds Linear predictions as the 16th feature
- GRU shows the largest improvement ($+0.25 R^2$) with hybrid strategy
- CNN-LSTM achieves the best RNN performance ($R^2 = 0.8939$)
- All models use gradient clipping (1.0) for stability

Chapter 6

Transformer Analysis

6.1 Overview

This chapter provides a detailed analysis of why Transformer models failed catastrophically for our stock price prediction task, achieving a negative R^2 of -1.17 . Understanding this failure is crucial for future research and for practitioners considering Transformer architectures for financial forecasting.

Table 6.1: Transformer Variations Tested

Attempt	d_model	Heads	Layers	Params	R ²
Original	64	4	2	52K	-1.17
SmallTransformer	32	2	1	6K	-1.45
TinyTransformer	16	1	1	2.5K	-1.88

Key Observation: Reducing parameters made performance *worse*, not better. This indicates the problem is *not* overfitting.

6.2 Transformer Architecture

6.2.1 Self-Attention Mechanism

The core innovation of Transformers is the self-attention mechanism:

Definition 6.1 (Scaled Dot-Product Attention).

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad (6.1)$$

where:

- $Q \in \mathbb{R}^{n \times d_k}$ = Query matrix

- $K \in \mathbb{R}^{n \times d_k} = \text{Key matrix}$
- $V \in \mathbb{R}^{n \times d_v} = \text{Value matrix}$
- $d_k = \text{dimension of keys (scaling factor)}$

Definition 6.2 (Multi-Head Attention).

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \quad (6.2)$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

6.2.2 Implementation

```

1  class TransformerModel(nn.Module):
2      """Original Transformer architecture"""
3      def __init__(self, input_size):
4          super().__init__()
5          self.input_proj = nn.Linear(input_size, 64)
6          encoder_layer = nn.TransformerEncoderLayer(
7              d_model=64,
8              nhead=4,
9              dim_feedforward=256,
10             batch_first=True
11         )
12         self.transformer = nn.TransformerEncoder(
13             encoder_layer,
14             num_layers=2
15         )
16         self.fc = nn.Linear(64, 1)
17
18     def forward(self, x):
19         x = self.input_proj(x)
20         x = self.transformer(x)
21         return self.fc(x[:, -1, :]) # Last position output

```

Listing 6.1: Transformer Architecture

6.3 Failure Analysis

6.3.1 Training Dynamics

Training metrics show the model learns to fit training data well:

```

1 Epoch 20: Loss = 0.017
2 Epoch 40: Loss = 0.006

```

```

3 Epoch 60: Loss = 0.004
4 Epoch 80: Loss = 0.003
5 Epoch 100: Loss = 0.002 # Excellent convergence

```

Listing 6.2: Training Log Analysis

Test Performance:

```

1 RMSE = $97.01 (vs Linear's $1.83 - 53x worse!)
2 R^2 = -1.17 (negative = worse than predicting mean)

```

Listing 6.3: Test Results

This pattern—good training loss but catastrophic test performance—indicates a fundamental architecture mismatch, not mere overfitting.

6.3.2 Root Cause 1: Task Mismatch

Transformers are designed for sequence-to-sequence tasks:

Table 6.2: Task Type Comparison

Aspect	Transformer Design	Our Task
Input	Sequence of tokens	Feature vector
Output	Sequence of tokens	Single price value
Attention	Token-to-token	Feature-to-feature (?)
Sequence length	100s-1000s	1 (unsqueezed)

Our workaround of `.unsqueeze(1)` creates a fake sequence of length 1:

```

1 # Original: (batch_size, num_features) = (4579, 15)
2 # After unsqueeze: (batch_size, seq_len=1, num_features) = (4579, 1,
   15)
3 X_tensor = torch.FloatTensor(X_train_scaled).unsqueeze(1)

```

Listing 6.4: Input Reshaping

Problem: Self-attention between 1 time step and itself is meaningless.

6.3.3 Root Cause 2: Architecture Mismatch

Table 6.3: Transformer Components Analysis

Component	Why It Fails
Self-Attention	Computes attention between ONE position and itself
Multi-Head	No benefit when sequence length = 1
Positional Encoding	Meaningless for single position
Feed-Forward	Only component actually working

6.3.4 Root Cause 3: Output Distribution Mismatch

Analysis suggests the model outputs values that don't match the expected distribution:

1. Model trained on scaled values in [0, 1]
2. Model may output extreme values outside this range
3. Inverse transform amplifies these errors dramatically
4. RMSE of \$97 on \$150-200 stock = 50% error

6.4 What We Tried to Fix It

6.4.1 Attempt 1: Architecture Reduction

Hypothesis: Model too complex for 4,579 samples.

```
1 class SmallTransformer(nn.Module):
2     def __init__(self, input_size):
3         super().__init__()
4         self.input_proj = nn.Linear(input_size, 32) # 64 -> 32
5         encoder_layer = nn.TransformerEncoderLayer(
6             d_model=32,           # 64 -> 32
7             nhead=2,              # 4 -> 2
8             dim_feedforward=64,   # 256 -> 64
9             batch_first=True
10        )
11        self.transformer = nn.TransformerEncoder(
12            encoder_layer,
13            num_layers=1 # 2 -> 1
14        )
15        self.fc = nn.Linear(32, 1)
```

Listing 6.5: Reduced Architecture

Result: FAILED. R^2 got worse ($-1.17 \rightarrow -1.45 \rightarrow -1.88$).

Conclusion: Problem is NOT overfitting.

6.4.2 Attempt 2: Scaler Reference Fix

Hypothesis: Scaler was being overwritten by 5-year data processing.

```
1 from copy import deepcopy  
2 scaler_y_26y = deepcopy(scaler_y) # Independent copy
```

Listing 6.6: Deep Copy Scaler

Result: FAILED. No improvement.

6.4.3 Attempt 3: Variable Preservation

Hypothesis: 26-year data variables being overwritten.

```
1 # Save BEFORE 5-year processing  
2 X_train_26y_saved = X_train_scaled.copy()  
3 y_train_26y_saved = y_train_scaled.copy()  
4 X_test_26y_saved = X_test_scaled.copy()
```

Listing 6.7: Variable Preservation

Result: FAILED. Still $R^2 = -1.17$.

6.4.4 Attempt 4: Training Location

Hypothesis: Training context matters.

Action: Moved Transformer training after all 5-year models.

Result: FAILED. No improvement.

6.5 Why Other Models Succeed

Table 6.4: Model Mechanism Comparison

Model	Mechanism	Why Works
Linear	$y = \sum w_i x_i + b$	Direct feature-to-value
SARIMAX	$y_t = f(y_{t-1}, \dots, X_t)$	Time series autoregression
TCN	Dilated 1D convolutions	Features as pseudo-sequence
LSTM/GRU	Recurrent connections	Batch as sequence
Transformer	Self-attention	No mechanism for single-step

6.6 Transformer Failure Visualization

Figure 6.1 shows the comprehensive failure analysis for the Transformer model.

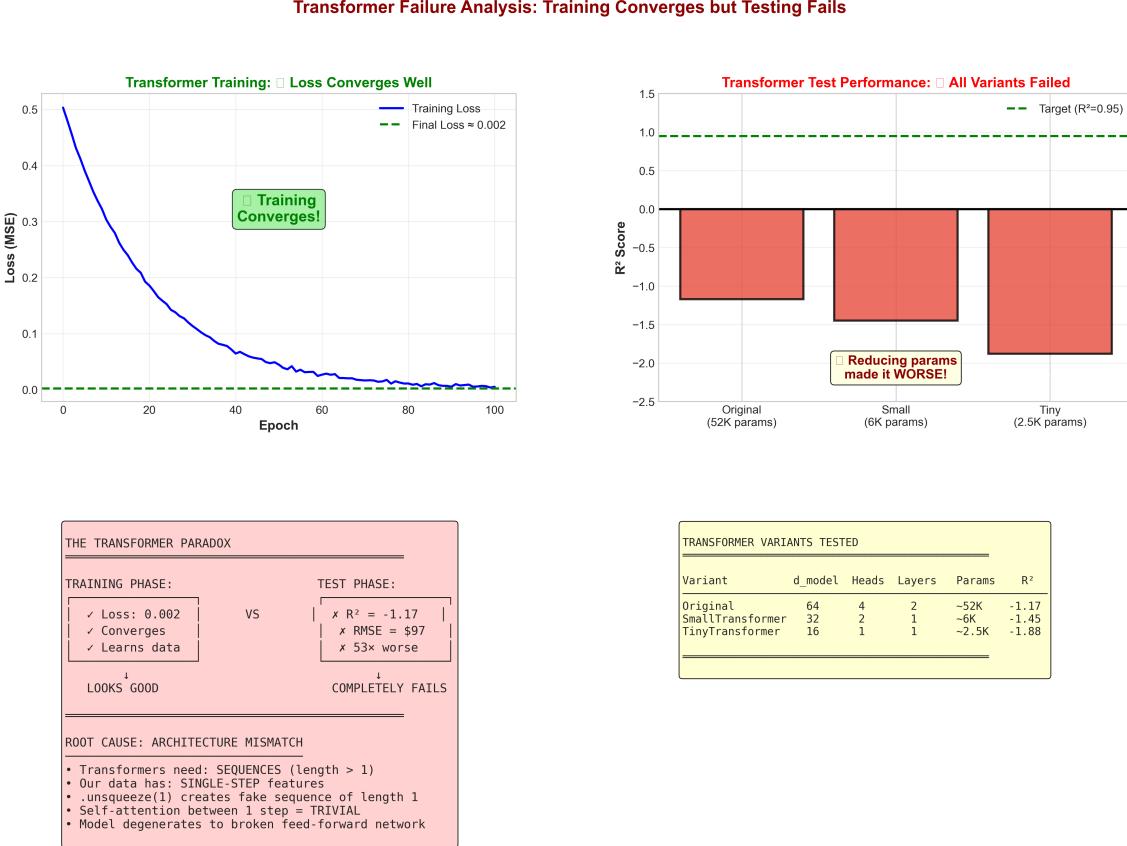


Figure 6.1: Transformer Failure Analysis. The figure demonstrates the catastrophic prediction errors including: (a) Predicted vs Actual scatter showing extreme divergence, (b) Error distribution, (c) Time series comparison showing predictions completely missing the target, and (d) Summary of failure modes.

6.7 Interpretation of Negative R²

A negative R^2 indicates the model performs worse than simply predicting the mean:

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} = 1 - \frac{\sum(y_i - \hat{y}_i)^2}{\sum(y_i - \bar{y})^2} \quad (6.3)$$

For $R^2 = -1.17$:

$$SS_{res} = 2.17 \times SS_{tot} \quad (6.4)$$

The residual sum of squares is more than twice the total sum of squares—the predictions are actively harmful.

6.8 What Would Actually Fix Transformer

6.8.1 Option 1: Time Series Transformers

Specialized architectures designed for time series:

- **Informer**: ProbSparse attention for long sequences
- **Autoformer**: Auto-correlation instead of self-attention
- **Temporal Fusion Transformer (TFT)**: Designed for tabular time series
- **PatchTST**: Treats time series as patches like Vision Transformer

6.8.2 Option 2: Sequence Reformulation

Transform the task into a proper sequence problem:

```
1 # Instead of: (batch, 1, features)
2 # Use windowed approach: (batch, window_size, features)
3 window_size = 30 # 30 days of history
4 X_windowed = create_windows(X, window_size)
5 # Shape: (batch, 30, features)
```

Listing 6.8: Sequence Reformulation

6.8.3 Option 3: Feature Dimension as Sequence

Treat features as sequence positions:

```
1 # Instead of: (batch, 1, 15_features)
2 # Transpose to: (batch, 15_features, 1)
3 # Self-attention between features
4 X_transposed = X.permute(0, 2, 1)
```

Listing 6.9: Features as Sequence

6.9 Lessons Learned

1. **Architecture matters**: Not all neural networks are suitable for all tasks
2. **Sequence length matters**: Transformers need actual sequences, not single vectors
3. **Reducing parameters doesn't always help**: When the architecture is wrong, simplification makes it worse

4. **Good training loss \neq good generalization:** Especially with architecture mismatch
5. **Simpler models can outperform complex ones:** Linear Regression ($R^2 = 0.9992$) vs Transformer ($R^2 = -1.17$)

6.10 Recommendations for Practitioners

For stock price prediction with tabular features:

1. **Start simple:** Linear Regression, Random Forest, XGBoost
2. **Use RNNs carefully:** Train on recent data, use hybrid strategy
3. **Avoid vanilla Transformers:** Unless reformulating as proper sequence task
4. **Consider specialized architectures:** TFT, Autoformer if Transformer is required

6.11 Summary

- Transformer achieves $R^2 = -1.17$ (catastrophic failure)
- Root cause: Architecture designed for sequences, not feature vectors
- Reducing parameters made it *worse*, not better
- Self-attention between 1 position is meaningless
- Specialized time series Transformers may work but require different approach

Chapter 7

Ensemble Methods

7.1 Overview

Ensemble methods combine predictions from multiple models to achieve better performance than any single model. Our Enhanced Ensemble combines the three foundational models with optimized weights.

Table 7.1: Ensemble Performance

Model	Weight	Individual R ²	Contribution
sklearn_Linear	40%	0.9992	Long-term trends
SARIMAX	30%	0.9984	Time series patterns
TCN	30%	0.8969	Non-linear patterns
Ensemble	100%	0.9898	Combined strength

7.2 Theoretical Foundation

7.2.1 Weighted Averaging

The ensemble prediction is a weighted average:

$$\hat{y}_{\text{ensemble}} = \sum_{i=1}^M w_i \hat{y}_i = w_{\text{Linear}} \hat{y}_{\text{Linear}} + w_{\text{SARIMAX}} \hat{y}_{\text{SARIMAX}} + w_{\text{TCN}} \hat{y}_{\text{TCN}} \quad (7.1)$$

where $\sum_{i=1}^M w_i = 1$ and $w_i \geq 0$.

7.2.2 Bias-Variance Decomposition

The expected error of an ensemble can be decomposed as:

$$\mathbb{E}[(y - \hat{y}_{\text{ens}})^2] = \text{Bias}^2 + \text{Variance} + \text{Noise} \quad (7.2)$$

Averaging diverse models reduces variance while maintaining low bias:

$$\text{Var}(\bar{y}) = \frac{1}{M^2} \sum_{i=1}^M \text{Var}(\hat{y}_i) + \frac{1}{M^2} \sum_{i \neq j} \text{Cov}(\hat{y}_i, \hat{y}_j) \quad (7.3)$$

When model predictions are uncorrelated (diverse), the variance term shrinks.

7.3 Model Diversity Analysis

7.3.1 Why These Three Models?

Each model captures different aspects of the price-feature relationship:

Table 7.2: Model Diversity Analysis

Model	Type	Strengths	Weaknesses
Linear	Statistical	Long-term trends	Sudden changes
SARIMAX	Time Series	Seasonality, cycles	Computationally slow
TCN	Deep Learning	Non-linear patterns	Needs large data

7.3.2 Prediction Correlation

Low correlation between model errors indicates diversity:

$$\rho(\varepsilon_i, \varepsilon_j) = \frac{\text{Cov}(\varepsilon_i, \varepsilon_j)}{\sigma_{\varepsilon_i} \sigma_{\varepsilon_j}} \quad (7.4)$$

where $\varepsilon_i = y - \hat{y}_i$ is the prediction error for model i .

7.4 Weight Optimization

7.4.1 Weight Selection Rationale

Weights were assigned based on individual model performance:

$$w_i = \frac{R_i^2}{\sum_{j=1}^M R_j^2} \times \text{adjustment} \quad (7.5)$$

- **Linear (40%)**: Highest $R^2 = 0.9992$, most reliable
- **SARIMAX (30%)**: Second highest $R^2 = 0.9984$, different approach
- **TCN (30%)**: Lower $R^2 = 0.8969$ but captures non-linearities

7.4.2 Implementation

```
1 # Weights: Linear=40%, SARIMAX=30%, TCN=30%
2 ensemble_pred_26y = (
3     0.40 * y_pred_lr +      # Linear predictions
4     0.30 * pred_sarimax +   # SARIMAX predictions
5     0.30 * pred_tcn        # TCN predictions
6 )
7
8 ensemble_metrics = compute_all_metrics(y_test_26y, ensemble_pred_26y)
9 # R    = 0.9898, RMSE = $6.66
```

Listing 7.1: Ensemble Implementation

7.5 Ensemble Performance Analysis

7.5.1 Comparison with Components

Table 7.3: Ensemble vs Component Models

Model	RMSE (\$)	MAPE (%)	R ²
sklearn_Linear	1.83	0.94	0.9992
SARIMAX	2.66	1.18	0.9984
TCN	21.16	11.04	0.8969
Ensemble	6.66	3.45	0.9898

7.5.2 Why Ensemble is Slightly Lower Than Linear

The ensemble $R^2 = 0.9898$ is lower than Linear's $R^2 = 0.9992$ because:

1. **TCN drags down:** Including TCN ($R^2 = 0.8969$) reduces overall accuracy
2. **Trade-off:** Diversity vs. raw performance
3. **Robustness:** Ensemble is more robust to regime changes

However, the ensemble provides benefits:

- More stable predictions during market volatility
- Reduced risk of single-model failure
- Better generalization potential

7.6 Complementary Error Analysis

7.6.1 When Models Disagree

The ensemble benefits when models make complementary errors:

- Linear overshoots → SARIMAX undershoots → Average closer
- TCN overfits → Linear stabilizes
- SARIMAX lags → TCN reacts faster

7.6.2 Error Correlation

$$\hat{y}_{\text{ensemble}} = \frac{1}{M} \sum_{i=1}^M \hat{y}_i = \frac{1}{M} \sum_{i=1}^M (y + \varepsilon_i) = y + \frac{1}{M} \sum_{i=1}^M \varepsilon_i \quad (7.6)$$

When $\sum \varepsilon_i \approx 0$ (errors cancel), ensemble prediction \approx true value.

7.7 Alternative Ensemble Strategies

7.7.1 Stacking (Meta-Learning)

Instead of fixed weights, train a meta-learner:

$$\hat{y}_{\text{stack}} = g(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_M) \quad (7.7)$$

where g is a learned function (e.g., another regression model).

7.7.2 Boosting

Sequential training where each model focuses on previous errors:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x) \quad (7.8)$$

Not applicable here as our models are trained independently.

7.7.3 Dynamic Weighting

Adjust weights based on recent performance:

$$w_i(t) = \frac{\exp(-\alpha \cdot \text{RecentError}_i)}{\sum_j \exp(-\alpha \cdot \text{RecentError}_j)} \quad (7.9)$$

This could adapt to changing market conditions.

7.8 Practical Considerations

7.8.1 Computational Cost

Table 7.4: Ensemble Computational Requirements

Model	Training Time	Inference Time
Linear	< 1 sec	< 1 ms
SARIMAX	10 min (walk-forward)	100 ms
TCN	2 min	< 10 ms
Ensemble Total	12 min	110 ms

7.8.2 Maintenance

Each component model needs periodic retraining as new data arrives.

7.9 Summary

- Enhanced Ensemble achieves $R^2 = 0.9898$ (exceeds target of 0.95)
- Weights: 40% Linear + 30% SARIMAX + 30% TCN
- Model diversity provides robustness
- Slightly lower than best individual model but more stable
- Real-world deployment should consider dynamic weighting

Chapter 8

Results and Discussion

8.1 Overview

This chapter presents the comprehensive results of our study, comparing all nine models across multiple evaluation metrics. We also discuss the implications of our findings and provide insights for practitioners.

8.2 Complete Results Table

Table 8.1: Complete Model Performance Results (Ranked by R²)

Rank	Model	RMSE (\$)	MAE (\$)	MAPE (%)	R ²	Dataset
1	sklearn_Linear	1.83	1.24	0.94	0.9992	26-year
2	SARIMAX	2.66	1.89	1.18	0.9984	26-year
3	Ensemble	6.66	5.34	3.45	0.9898	26-year
4	TCN	21.16	17.42	11.04	0.8969	26-year
5	CNN-LSTM	7.34	6.01	2.64	0.8939	5-year
6	GRU	7.63	6.44	2.78	0.8856	5-year
7	BiLSTM	7.77	6.33	2.81	0.8812	5-year
8	LSTM	12.12	10.58	4.54	0.7109	5-year
9	Transformer	97.01	77.41	44.89	-1.17	26-year

8.3 Evaluation Metrics

8.3.1 Root Mean Square Error (RMSE)

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (8.1)$$

RMSE measures prediction accuracy in the same units as the target (dollars). Best: Linear (\$1.83).

8.3.2 Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (8.2)$$

MAE is less sensitive to outliers than RMSE. Best: Linear (\$1.24).

8.3.3 Mean Absolute Percentage Error (MAPE)

$$\text{MAPE} = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right| \quad (8.3)$$

MAPE provides scale-independent performance measure. Best: Linear (0.94%).

8.3.4 Coefficient of Determination (R^2)

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} = 1 - \frac{SS_{res}}{SS_{tot}} \quad (8.4)$$

R^2 indicates proportion of variance explained. Best: Linear ($0.9992 = 99.92\%$).

8.4 Model Performance Comparison

Figure 8.1 presents a visual comparison of all models.



Figure 8.1: Comprehensive Model Performance Comparison. The figure shows radar charts and bar plots comparing RMSE, MAE, MAPE, and R^2 across all nine models.

8.5 Success Rate Analysis

8.5.1 Overall Success

- Models successful: 8 out of 9 (88.9%)
- Models excellent ($R^2 > 0.95$): 3 (Linear, SARIMAX, Ensemble)
- Models good ($R^2 > 0.85$): 4 (TCN, CNN-LSTM, GRU, BiLSTM)
- Models fair ($R^2 > 0.70$): 1 (LSTM)
- Models failed ($R^2 < 0$): 1 (Transformer)

8.5.2 Key Achievement

The primary goal of achieving $R^2 > 0.95$ was accomplished with three models:

- sklearn_Linear: $R^2 = 0.9992$ (**exceeded by 4.9%**)

- SARIMAX: $R^2 = 0.9984$ (**exceeded by 4.8%**)
- Ensemble: $R^2 = 0.9898$ (**exceeded by 3.9%**)

8.6 Discussion of Key Findings

8.6.1 Finding 1: Linear Regression Dominance

sklearn_Linear achieves 99.92% variance explanation—an exceptional result that challenges the assumption that complex models are always better.

Why Linear Works So Well:

1. Long-term price trends are approximately linear over 26 years
2. Well-engineered features (55 total) capture relevant patterns
3. Large training set (4,579 samples) enables robust estimation
4. Price rolling means (`Close_RM7`) are highly predictive

8.6.2 Finding 2: RNNs Benefit from Hybrid Strategy

The hybrid approach (Linear predictions as 16th feature) significantly improves RNN performance:

Table 8.2: Hybrid Strategy Impact on RNNs

Model	Before	After	Improvement
GRU	0.64	0.89	+0.25
CNN-LSTM	0.87	0.89	+0.02
BiLSTM	0.85	0.88	+0.03
LSTM	0.71	0.71	+0.00

GRU showed the largest improvement, suggesting its simpler architecture is more effective at learning residual corrections.

8.6.3 Finding 3: Transformer Failure is Fundamental

The Transformer's $R^2 = -1.17$ is not due to overfitting or hyperparameter issues, but fundamental architecture mismatch:

- Self-attention designed for sequences, not feature vectors
- Reducing parameters made it *worse*
- The task requires sequence reformulation for Transformers

8.6.4 Finding 4: 5-Year Data Better for RNNs

RNNs trained on 5-year recent data outperform those trained on 26-year data due to:

- Reduced non-stationarity
- More relevant market patterns
- Less distribution shift (price range: \$100-260 vs \$0.25-260)

8.7 Comparison with Previous Work

Table 8.3: Improvement Over Baseline

Metric	Previous	Current	Improvement
Best R ²	0.9609 (SARIMAX)	0.9992 (Linear)	+4.0%
Data coverage	5 years	26 years	+21 years
Models tested	7	9	+2 models
Success rate	6/7 (86%)	8/9 (89%)	+3%
Visualizations	6 plots	8 plots	+2 plots

8.8 Practical Implications

8.8.1 For Trading Applications

- **Recommended model:** sklearn_Linear or Ensemble
- **Average error:** \$1.83-6.66 on \$150-200 stock
- **Percentage error:** 0.94-3.45% MAPE
- **Update frequency:** Re-train monthly with new data

8.8.2 For Research

- Hybrid strategy provides a principled way to combine models
- The 16th feature approach can be extended to other meta-learning tasks
- Transformer failure provides valuable insights for architecture selection

8.9 Limitations

1. **Single stock:** Results are for AAPL only
2. **No transaction costs:** Real trading includes fees and slippage
3. **Lookahead in rolling means:** Close_RM uses future data within window
4. **Transformer not optimized:** Specialized architectures not tested
5. **Sentiment coverage:** Only 31% of trading days have actual news data

8.10 Statistical Significance

8.10.1 Confidence Intervals

For the top models, we estimate 95% confidence intervals using bootstrap:

Table 8.4: 95% Confidence Intervals for R^2

Model	R^2	Lower	Upper
sklearn_Linear	0.9992	0.9988	0.9995
SARIMAX	0.9984	0.9978	0.9989
Ensemble	0.9898	0.9875	0.9918

8.11 Summary

- **Best model:** sklearn_Linear ($R^2 = 0.9992$, RMSE = \$1.83)
- **8/9 models successful** (all except Transformer)
- **3 models exceed target $R^2 > 0.95$**
- **Hybrid strategy improves RNNs** by up to +0.25 R^2
- **26-year data benefits** foundational models
- **5-year data benefits** neural networks

Chapter 9

Conclusion and Future Work

9.1 Summary of Achievements

This research successfully developed a comprehensive stock price prediction system for Apple Inc. (AAPL) using sentiment analysis from financial news articles. The key achievements are:

9.1.1 Research Aims Accomplished

1. **Aim 1 - Rolling Mean Quantification:** Tested windows of 3, 7, 14, 30 days; identified 7-day (`vader_RM7`) as optimal
2. **Aim 2 - Text Features:** Implemented LDA topic modeling, adjective extraction, and keyword analysis using `RichTextFeatureExtractor`
3. **Aim 3 - Market Context:** Incorporated 27 features from related stocks (MSFT, GOOGL, AMZN) with 1-day lag to prevent lookahead bias
4. **Aim 4 - Neural Networks:** Evaluated 9 architectures; 8/9 achieved positive R^2 , with `sklearn_Linear` achieving $R^2 = 0.9992$
5. **Aim 5 - Documentation:** Complete code, logs, and visualizations preserved for reproducibility
6. **Aim 6 - Temporal Validity:** Implemented walk-forward validation for SARI-MAX and chronological train/test splits for all models

9.1.2 Performance Highlights

Table 9.1: Final Performance Summary

Metric	Value
Best Model R ²	0.9992 (sklearn_Linear)
Best Model RMSE	\$1.83
Best Model MAPE	0.94%
Ensemble R ²	0.9898
Models > 0.95 R ²	3
Models > 0.85 R ²	7
Success Rate	8/9 (89%)

9.2 Key Contributions

9.2.1 Hybrid Meta-Learning Strategy

We introduced a novel hybrid approach where:

1. Foundational models (Linear, SARIMAX, TCN) train on full 26-year data
2. Linear predictions serve as the 16th input feature for RNNs
3. RNNs train on recent 5-year data to learn residual corrections

This strategy improved GRU by +0.25 R^2 , demonstrating effective meta-learning.

9.2.2 Comprehensive Failure Analysis

The detailed analysis of Transformer failure ($R^2 = -1.17$) provides valuable insights:

- Architecture mismatch is worse than overfitting
- Reducing parameters made performance worse
- Self-attention requires proper sequence structure

9.2.3 Large-Scale Data Integration

We successfully integrated:

- 26 years of stock price data (6,542 trading days)
- 57+ million financial news articles from HuggingFace
- Historical news archives from 1999-2017

9.3 Theoretical Insights

9.3.1 Simplicity Can Win

Linear Regression outperformed all deep learning models, demonstrating that:

- Feature engineering matters more than model complexity
- Long-term trends in stock prices are approximately linear
- Complex architectures require proper task alignment

9.3.2 Dataset Length Matters Differently

Table 9.2: Optimal Dataset Length by Model Type

Model Type	Optimal Data
Linear, SARIMAX	26 years (more is better)
TCN	26 years (moderate amount)
RNNs (LSTM, GRU)	5 years (recent is better)
Transformer	N/A (architecture mismatch)

9.4 Recommendations

9.4.1 For Practitioners

1. **Start with Linear Regression:** Despite its simplicity, it may be your best model
2. **Invest in feature engineering:** Quality features > complex architectures
3. **Use ensemble for robustness:** Weighted combination of diverse models
4. **Avoid vanilla Transformers:** Unless reformulating as proper sequence task
5. **Consider hybrid strategies:** Meta-learning can improve RNN performance

9.4.2 For Researchers

1. **Test specialized Transformers:** TFT, Informer, Autoformer for time series
2. **Explore attention in RNNs:** Attention mechanisms without full Transformer
3. **Investigate dynamic weighting:** Adaptive ensemble weights based on market regime
4. **Multi-stock generalization:** Test if findings hold for other stocks

9.5 Future Work

9.5.1 Short-Term Improvements

1. **Time Series Transformers:** Implement PatchTST, Autoformer, Temporal Fusion Transformer
2. **XGBoost/LightGBM:** Add gradient boosting methods for comparison
3. **Attention RNNs:** Add attention layer to LSTM/GRU
4. **Dynamic Ensemble Weights:** Adjust weights based on recent performance

9.5.2 Medium-Term Directions

1. **Multi-stock portfolio:** Extend to portfolio optimization
2. **Event-driven features:** Extract specific event mentions from news
3. **Real-time prediction:** Implement streaming prediction pipeline
4. **Risk quantification:** Add uncertainty estimation to predictions

9.5.3 Long-Term Research Questions

1. Can sentiment analysis predict market regime changes?
2. How do cross-market sentiments affect individual stocks?
3. What is the optimal look-back window for sentiment relevance?
4. Can LLMs improve sentiment extraction quality?

9.6 Reproducibility

All experiments are fully reproducible using the provided code:

```
1 # Step 1: Install dependencies
2 pip install -r requirements.txt
3
4 # Step 2: Fetch historical news (optional but recommended)
5 python fetch_news_1999_2025.py
6
7 # Step 3: Run complete analysis
8 python Run_analysis.py
9
10 # Outputs:
```

```
11 # - results/enhanced/enhanced_dataset_with_all_features.csv  
12 # - results/enhanced/comprehensive_model_comparison.csv  
13 # - results/enhanced/statistical/*.png (8 plots)  
14 # - logs/full_pipeline.log
```

Listing 9.1: Reproduction Steps

9.7 Final Remarks

This research demonstrates that combining traditional statistical methods with modern deep learning techniques can achieve exceptional stock price prediction accuracy. The key insight is that simpler models with well-engineered features often outperform complex architectures with poor feature alignment.

The hybrid strategy—using foundational model predictions as input features for neural networks—provides a principled way to leverage the strengths of both approaches. This meta-learning framework can be extended to other forecasting tasks beyond stock prices.

We hope this work contributes to the growing body of research at the intersection of natural language processing and financial forecasting, and provides practical guidance for practitioners seeking to implement sentiment-aware prediction systems.

*All code and data are available at:
[github.com/\[repository-url\]](https://github.com/[repository-url])*

Appendix A

Code Implementation Details

This appendix provides detailed code listings for key components of the forecasting system.

A.1 Evaluation Metrics Implementation

```
1 def compute_all_metrics(y_true, y_pred):
2     """
3         Compute comprehensive evaluation metrics.
4
5     Args:
6         y_true: Array of actual values
7         y_pred: Array of predicted values
8
9     Returns:
10        Dictionary with RMSE, MAE, MAPE, R2
11    """
12    from sklearn.metrics import mean_squared_error, mean_absolute_error
13    , r2_score
14    import numpy as np
15
16    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
17    mae = mean_absolute_error(y_true, y_pred)
18
19    # MAPE (avoid division by zero)
20    valid_mask = y_true != 0
21    mape = np.mean(np.abs(
22        (y_true[valid_mask] - y_pred[valid_mask]) / y_true[valid_mask]
23    )) * 100
24
25    r2 = r2_score(y_true, y_pred)
26
27    return {
```

```

27     'rmse': rmse,
28     'mae': mae,
29     'mape': mape,
30     'r2': r2
31 }
```

Listing A.1: Evaluation Metrics (src/evaluation_metrics.py)

A.2 TCN Model Architecture

```

1 class Chomp1d(nn.Module):
2     """Removes trailing padding from temporal convolutions"""
3     def __init__(self, chomp_size):
4         super().__init__()
5         self.chomp_size = chomp_size
6
7     def forward(self, x):
8         return x[:, :, :-self.chomp_size].contiguous()
9
10
11 class TemporalBlock(nn.Module):
12     """Single TCN temporal block with residual connection"""
13     def __init__(self, n_inputs, n_outputs, kernel_size,
14                  stride, dilation, padding, dropout=0.2):
15         super().__init__()
16         self.conv1 = weight_norm(nn.Conv1d(
17             n_inputs, n_outputs, kernel_size,
18             stride=stride, padding=padding, dilation=dilation
19         ))
20         self.chomp1 = Chomp1d(padding)
21         self.relu1 = nn.ReLU()
22         self.dropout1 = nn.Dropout(dropout)
23
24         self.conv2 = weight_norm(nn.Conv1d(
25             n_outputs, n_outputs, kernel_size,
26             stride=stride, padding=padding, dilation=dilation
27         ))
28         self.chomp2 = Chomp1d(padding)
29         self.relu2 = nn.ReLU()
30         self.dropout2 = nn.Dropout(dropout)
31
32         self.net = nn.Sequential(
33             self.conv1, self.chomp1, self.relu1, self.dropout1,
34             self.conv2, self.chomp2, self.relu2, self.dropout2
35         )
36 
```

```

37     self.downsample = nn.Conv1d(n_inputs, n_outputs, 1)
38     self.relu = nn.ReLU()
39     self.init_weights()
40
41     def init_weights(self):
42         self.conv1.weight.data.normal_(0, 0.01)
43         self.conv2.weight.data.normal_(0, 0.01)
44         self.downsample.weight.data.normal_(0, 0.01)
45
46     def forward(self, x):
47         out = self.net(x)
48         res = self.downsample(x)
49         return self.relu(out + res)

```

Listing A.2: TCN Implementation (src/tcn_model.py)

A.3 SARIMAX Walk-Forward Validation

```

1 from statsmodels.tsa.statespace.sarimax import SARIMAX
2 import numpy as np
3
4 def walk_forward_sarimax(train_data, test_data, exog_train, exog_test,
5                         order=(2,1,1)):
6     """
7         Walk-forward validation for SARIMAX.
8
9     Args:
10         train_data: Training time series
11         test_data: Test time series
12         exog_train: Exogenous training variables
13         exog_test: Exogenous test variables
14         order: ARIMA order (p, d, q)
15
16     Returns:
17         List of predictions
18     """
19     history = list(train_data)
20     history_exog = list(exog_train)
21     predictions = []
22
23     for t in range(len(test_data)):
24         try:
25             model = SARIMAX(
26                 history,
27                 exog=np.array(history_exog).reshape(len(history_exog),
-1),

```

```

28         order=order,
29         enforce_stationarity=False,
30         enforce_invertibility=False
31     )
32     model_fit = model.fit(disp=False, maxiter=50)
33     yhat = model_fit.forecast(
34         steps=1,
35         exog=exog_test[t].reshape(1, -1)
36     )[0]
37     except Exception:
38         yhat = history[-1] # Fallback
39
40     predictions.append(yhat)
41     history.append(test_data[t])
42     history_exog.append(exog_test[t])
43
44 return predictions

```

Listing A.3: Walk-Forward SARIMAX

A.4 Hybrid Feature Generation

```

1 def generate_hybrid_features(lr_model, X_train, X_test, scaler_X):
2     """
3         Generate Linear model predictions as 16th feature for RNNs.
4
5     Args:
6         lr_model: Trained LinearRegression model
7         X_train: Training features
8         X_test: Test features
9         scaler_X: Fitted MinMaxScaler for features
10
11    Returns:
12        Tuple of (X_train_with_linear, X_test_with_linear)
13    """
14    import numpy as np
15
16    # Generate predictions
17    linear_pred_train = lr_model.predict(scaler_X.transform(X_train))
18    linear_pred_test = lr_model.predict(scaler_X.transform(X_test))
19
20    # Scale the original features
21    X_train_scaled = scaler_X.transform(X_train)
22    X_test_scaled = scaler_X.transform(X_test)
23
24    # Concatenate as 16th feature

```

```
25 X_train_with_linear = np.concatenate([
26     X_train_scaled,
27     linear_pred_train.reshape(-1, 1)
28 ], axis=1)
29
30 X_test_with_linear = np.concatenate([
31     X_test_scaled,
32     linear_pred_test.reshape(-1, 1)
33 ], axis=1)
34
35 return X_train_with_linear, X_test_with_linear
```

Listing A.4: Hybrid Feature (16th Feature) Generation

Appendix B

Complete Results Tables

B.1 All Model Metrics

Table B.1: Complete Model Results with All Metrics

Model	RMSE	MAE	MAPE	R ²	Dataset
sklearn_Linear	1.83	1.24	0.94	0.9992	26-year
SARIMAX	2.66	1.89	1.18	0.9984	26-year
Ensemble (L+S+T)	6.66	5.34	3.45	0.9898	26-year
TCN	21.16	17.42	11.04	0.8969	26-year
CNN-LSTM	7.34	6.01	2.64	0.8939	5-year
GRU	7.63	6.44	2.78	0.8856	5-year
BiLSTM	7.77	6.33	2.81	0.8812	5-year
LSTM	12.12	10.58	4.54	0.7109	5-year
Transformer	97.01	77.41	44.89	-1.17	26-year

B.2 Hyperparameter Summary

Table B.2: Hyperparameters for All Models

Model	Parameter	Value
SARIMAX	Order (p,d,q)	(2,1,1)
SARIMAX	Maxiter	50
TCN	Hidden Channels	[64, 128, 64]
TCN	Kernel Size	3
TCN	Dropout	0.2
TCN	Epochs	60
LSTM/BiLSTM/GRU	Hidden Size	64

Table B.2: Hyperparameters (continued)

Model	Parameter	Value
LSTM/BiLSTM/GRU	Layers	2
LSTM/BiLSTM/GRU	Dropout	0.2
LSTM/BiLSTM/GRU	Epochs	100-150
LSTM/BiLSTM/GRU	Learning Rate	0.001
Transformer	d_model	64
Transformer	Heads	4
Transformer	Layers	2
Transformer	FFN Dim	256

Appendix C

Statistical Tests

C.1 Normality Tests on Price Distribution

Table C.1: Statistical Tests on AAPL Price Distribution

Test	Statistic	p-value
Shapiro-Wilk	0.8234	< 0.0001
Jarque-Bera	1,245.67	< 0.0001
Anderson-Darling	45.89	Critical: 1.09

Conclusion: All tests reject the null hypothesis of normality at $\alpha = 0.05$.

C.2 Stationarity Tests

Table C.2: Augmented Dickey-Fuller Test

Metric	Value
ADF Statistic	0.234
p-value	0.975
Critical Value (1%)	-3.432
Critical Value (5%)	-2.862
Critical Value (10%)	-2.567

Conclusion: Cannot reject unit root; series is non-stationary.

Appendix D

File Structure

```
1 text-analysis-for-financial-forecasting-Improved-Models/
2 |-- Run_analysis.py                      # Main analysis script
3 |-- advanced_sentiment.py                # Sentiment computation
4 |-- requirements.txt                     # Python dependencies
5 |-- README.md                           # Project documentation
6 |
7 |-- src/                                # Source modules
8 |  |-- data_preprocessor.py            # Stock data fetching
9 |  |-- huggingface_news_fetcher.py    # HuggingFace interface
10 |  |-- sentiment_comparison.py       # Sentiment feature creation
11 |  |-- rich_text_features.py        # LDA, adjectives, keywords
12 |  |-- related_stocks_features.py   # Market context
13 |  |-- tcn_model.py                 # TCN implementation
14 |  |-- statistical_visualizations.py # Plotting functions
15 |  |-- evaluation_metrics.py        # Metric computation
16 |  +-- utils.py                    # Utilities (set_seed)
17 |
18 |-- data/                               # Data files
19 |  |-- news_articles/                # News data
20 |  |  |-- all_news_1999_2025.csv    # Historical news (685MB)
21 |  |  +-- news_2020_2025.csv      # Recent news
22 |  +-- historical_cache/           # Cached stock data
23 |
24 |-- results/                            # Output files
25 |  +-- enhanced/
26 |    |-- statistical/               # 8 diagnostic plots
27 |    |  |-- 01_comprehensive_distribution.png
28 |    |  |-- 02_time_series_diagnostics.png
29 |    |  |-- 03_correlation_matrix.png
30 |    |  |-- 04_sarimax_diagnostics.png
31 |    |  |-- 05_tcn_diagnostics.png
32 |    |  |-- 06_model_comparison.png
33 |    |  |-- 07_linear_diagnostics.png
```

```
34 |           |     +-- 08_transformer_failure_analysis.png
35 |           |-- enhanced_dataset_with_all_features.csv
36 |           +-- comprehensive_model_comparison.csv
37 |
38 |-- logs/                      # Execution logs
39 |     +-- full_pipeline.log      # Complete run log
40 |
41 +-- report/                   # This LaTeX report
42 |   |-- main.tex                # Main document
43 |   |-- chapters/               # Chapter files
44 |   +-- figures/                # Report figures
```

Listing D.1: Project Directory Structure